
Computer Architecture

Ngo Lam Trung, Pham Ngoc Hung, Hoang Van Hiep
Department of Computer Engineering
School of Information and Communication Technology (SoICT)
Hanoi University of Science and Technology
E-mail: [trungnl, hungpn, hiephv]@soict.hust.edu.vn

Chapter 4: Arithmetic for Computers

[with materials from *COD, RISC-V 2nd Edition*, Patterson & Hennessy 2021,
and M.J. Irwin's presentation, PSU 2008,
The RISC-V Instruction Set Manual, Volume I, ver. 2.2]

What are stored inside computer?

- ❑ Data, of course!
 - ❑ Audio, video, image, drawings,...
 - ❑ Documents, personal information,...
 - ❑ Finance record, corporate business data,...
 - ❑ ...
- ❑ Complex data are constructed from basic data types.
 - ❑ Integers
 - ❑ Real numbers (Floating point)
 - ❑ Symbols (Characters)
- ❑ All are represented as binary numbers.

Content

- ❑ (Super) Basics of logic design
- ❑ Integer representation
- ❑ Integer arithmetic (inside computer)
- ❑ Floating point number representation and arithmetic

Unsigned Binary Integers

- ❑ Using n-bit binary number to represent non-negative integer

$$\begin{aligned} X &= X_{n-1}X_{n-2}\dots X_1X_0 \\ &= X_{n-1}2^{n-1} + X_{n-2}2^{n-2} + \dots + X_12^1 + X_02^0 \end{aligned}$$

- ❑ Range: 0 to $+2^n - 1$

- ❑ Example

$$\begin{aligned} &0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 0000\ 1011_2 \\ &= 0 + \dots + 1 \times 2^3 + 0 \times 2^2 + 1 \times 2^1 + 1 \times 2^0 \\ &= 0 + \dots + 8 + 0 + 2 + 1 = 11_{10} \end{aligned}$$

- ❑ Data range using 32 bits

$$0 \text{ to } 2^{32}-1 = 4,294,967,295$$

Eg: 32 bit Unsigned Binary Integers

Hex	Binary	Decimal
0x00000000	0...0000	0
0x00000001	0...0001	1
0x00000002	0...0010	2
0x00000003	0...0011	3
0x00000004	0...0100	4
0x00000005	0...0101	5
0x00000006	0...0110	6
0x00000007	0...0111	7
0x00000008	0...1000	8
0x00000009	0...1001	9
	...	
0xFFFFFFF0	1...1111	$2^{32}-1$
0xFFFFFFF1	1...1110	$2^{32}-2$
0xFFFFFFF2	1...1101	$2^{32}-3$
0xFFFFFFF3	1...1100	$2^{32}-4$

Exercise

- ❑ Convert from decimal to 32-bit binary integers

25 = 0000 0000 0000 0000 0000 0000 0001 1001

125 = 0000 0000 0000 0000 0000 0000 0111 1101

255 = 0000 0000 0000 0000 0000 0000 1111 1111

- ❑ Convert 32-bit binary integers to decimal

0000 0000 0000 0000 0000 0000 1100 1111 = 207

0000 0000 0000 0000 0000 0001 0011 0011 = 307

Signed binary integers

- Using n-bit binary number to represent integer, including negative values

$$\begin{aligned} X &= X_{n-1}X_{n-2}\dots X_1X_0 \\ &= -X_{n-1}2^{n-1} + X_{n-2}2^{n-2} + \dots + X_12^1 + X_02^0 \end{aligned}$$

- Range: -2^{n-1} to $+2^{n-1} - 1$**

- Example**

$$\begin{aligned} &1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1111\ 1100_2 \\ &= -1 \times 2^{31} + 1 \times 2^{30} + \dots + 1 \times 2^2 + 0 \times 2^1 + 0 \times 2^0 \\ &= -2,147,483,648 + 2,147,483,644 = -4_{10} \end{aligned}$$

- Using 32 bits**

$$-2,147,483,648 \text{ to } +2,147,483,647$$

Signed integer negation

- Given $x = x_{n-1}x_{n-2} \dots x_1x_0$, how to calculate $-x$?
- Let \bar{x} = 1's complement of x

$$\bar{x} = 1111 \dots 11_2 - x$$

$$(1 \rightarrow 0, 0 \rightarrow 1)$$

Then

$$\bar{x} + x = 1111 \dots 11_2 = -1$$

$$\rightarrow \bar{x} + 1 = -x$$

- **Example: find binary representation of -2**

$$+2 = 0000 \ 0000 \dots 0010_2$$

$$\begin{aligned} -2 &= 1111 \ 1111 \dots 1101_2 + 1 \\ &= 1111 \ 1111 \dots 1110_2 \end{aligned}$$

Exercise

□ Find 16 bit signed integer representation of

$$16 = 0000\ 0000\ 0001\ 0000$$

$$-16 = 1111\ 1111\ 1111\ 0000$$

$$100 = 0000\ 0000\ 0110\ 0100$$

$$-100 = 1111\ 1111\ 1001\ 1100$$

Sign extension

- ❑ Given n-bit integer $x = x_{n-1}x_{n-2} \dots x_1x_0$
- ❑ Find corresponding m-bit representation ($m > n$) with the same numeric value

$$x = x_{m-1}x_{m-2} \dots x_1x_0$$

- ❑ → Replicate the sign bit to the left

- ❑ Examples: 8-bit to 16-bit

+2: 0000 0010 => 0000 0000 0000 0010

-2: 1111 1110 => 1111 1111 1111 1110

Instruction to work with sign/unsigned

- ❑ lb/lbu, lh/lhu
- ❑ blt/bltu, bge/bgeu
- ❑ slt/sltu, slti/sltiu
- ❑ div/divu, rem/remu

Example

- ❑ What is the output of the following program?
- ❑ What if the blt instruction is replaced by bltu?

```
        li t0, 20
        li t1, -20
        blt t1, t0, else
        li a0, 1
        j print
else:
        li a0, 0
print:
        li a7, 1
        ecall
```

Example

❑ What are the decimal values of s0, s1, s2, s3?

```
.data
    x: .byte 20
    y: .byte -20
.text
    la t0, x
    la t1, y
    lb  s0, 0(t0)
    lbu s1, 0(t0)
    lb  s2, 0(t1)
    lbu s3, 0(t1)
```

Addition and subtraction

❑ Addition

- ❑ Similar to what you do to add two numbers manually
- ❑ Digits are added bit by bit from right to left
- ❑ Carries passed to the next digit to the left

❑ Subtraction

- ❑ Negate the second operand then add to the first operand

$$\begin{array}{r} + \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0111_{\text{two}} = 7_{\text{ten}} \\ \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0110_{\text{two}} = 6_{\text{ten}} \\ \hline \quad 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 0000 \ 1101_{\text{two}} = 13_{\text{ten}} \end{array}$$

Carryout and Overflow

- ❑ Carryout: adding or subtracting n-bit binary numbers result in carryout to or borrow from bit n+1.
- ❑ Overflow: adding or subtracting n-bit signed integers result in a value that cannot be represented by a n-bit signed integer.
 - ❑ When adding operands with different signs or when subtracting operands with the same sign, overflow can *never* occur

Operation	Operand A	Operand B	Result indicating overflow
A + B	≥ 0	≥ 0	< 0
A + B	< 0	< 0	≥ 0
A - B	≥ 0	< 0	< 0
A - B	< 0	≥ 0	≥ 0

Examples

- ❑ All numbers are 8-bit signed integer

$$12 + 8 =$$

$$122 + 8 =$$

$$122 + 80 =$$

$$122 - 80 =$$

Basics of logic design (Appendix A)

- ❑ Boolean logic: logic variable and operators
- ❑ Logic variable: values of 1 (TRUE) or 0 (FALSE)
- ❑ Basic operators: AND, OR, NOT
 - ❑ A AND B : $A \cdot B$ hay AB
 - ❑ A OR B : $A + B$
 - ❑ NOT A : \overline{A}
 - ❑ Order: NOT > AND > OR
- ❑ Additional operators: NAND, NOR, XOR
 - ❑ A NAND B: $\overline{A \cdot B}$
 - ❑ A NOR B : $\overline{A + B}$
 - ❑ A XOR B: $A \oplus B = A \bullet \overline{B} + \overline{A} \bullet B$

Truth tables

A	B	A AND B $A \bullet B$
0	0	0
0	1	0
1	0	0
1	1	1

A	B	A OR B $A + B$
0	0	0
0	1	1
1	0	1
1	1	1

A	NOT A $\neg A$
0	1
1	0


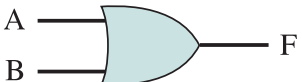
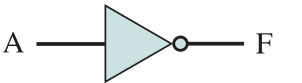

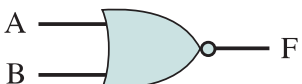
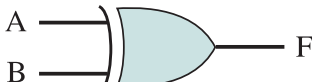
Unary operator NOT

A	B	A NAND B $A \bullet B$
0	0	1
0	1	1
1	0	1
1	1	0

A	B	A XOR B $A \oplus B$
0	0	0
0	1	1
1	0	1
1	1	0

A	B	A NOR B $A + B$
0	0	1
0	1	0
1	0	0
1	1	0

Logic gates

Name	Graphical Symbol	Algebraic Function	Truth Table															
AND		$F = A \bullet B$ or $F = AB$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	F	0	0	0	0	1	0	1	0	0	1	1	1
A	B	F																
0	0	0																
0	1	0																
1	0	0																
1	1	1																
OR		$F = A + B$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>1</td></tr></table>	A	B	F	0	0	0	0	1	1	1	0	1	1	1	1
A	B	F																
0	0	0																
0	1	1																
1	0	1																
1	1	1																
NOT		$F = \overline{A}$ or $F = A'$	<table><tr><th>A</th><th>F</th></tr><tr><td>0</td><td>1</td></tr><tr><td>1</td><td>0</td></tr></table>	A	F	0	1	1	0									
A	F																	
0	1																	
1	0																	
NAND		$F = \overline{AB}$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	1	0	1	1	1	0	1	1	1	0
A	B	F																
0	0	1																
0	1	1																
1	0	1																
1	1	0																
NOR		$F = \overline{A + B}$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>1</td></tr><tr><td>0</td><td>1</td><td>0</td></tr><tr><td>1</td><td>0</td><td>0</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	1	0	1	0	1	0	0	1	1	0
A	B	F																
0	0	1																
0	1	0																
1	0	0																
1	1	0																
XOR		$F = A \oplus B$	<table><tr><th>A</th><th>B</th><th>F</th></tr><tr><td>0</td><td>0</td><td>0</td></tr><tr><td>0</td><td>1</td><td>1</td></tr><tr><td>1</td><td>0</td><td>1</td></tr><tr><td>1</td><td>1</td><td>0</td></tr></table>	A	B	F	0	0	0	0	1	1	1	0	1	1	1	0
A	B	F																
0	0	0																
0	1	1																
1	0	1																
1	1	0																

Laws of Boolean algebra

$$A \cdot B = B \cdot A$$

$$A \cdot (B + C) = (A \cdot B) + (A \cdot C)$$

$$1 \cdot A = A$$

$$A \cdot \bar{A} = 0$$

$$0 \cdot A = 0$$

$$A \cdot A = A$$

$$A \cdot (B \cdot C) = (A \cdot B) \cdot C$$

$$\overline{A \cdot B} = \bar{A} + \bar{B} \text{ (DeMorgan's law)}$$

$$A + B = B + A$$

$$A + (B \cdot C) = (A + B) \cdot (A + C)$$

$$0 + A = A$$

$$A + \bar{A} = 1$$

$$1 + A = 1$$



$$A + A = A$$

$$A + (B + C) = (A + B) + C$$

$$\overline{A + B} = \bar{A} \cdot \bar{B} \text{ (DeMorgan's law)}$$


Example: multiplexor

- ❑ Depending on S , output C is equal to one of the two inputs A , B
- ❑ Explain how this circuit works?



Adder implementation

❑ 1-bit full adder



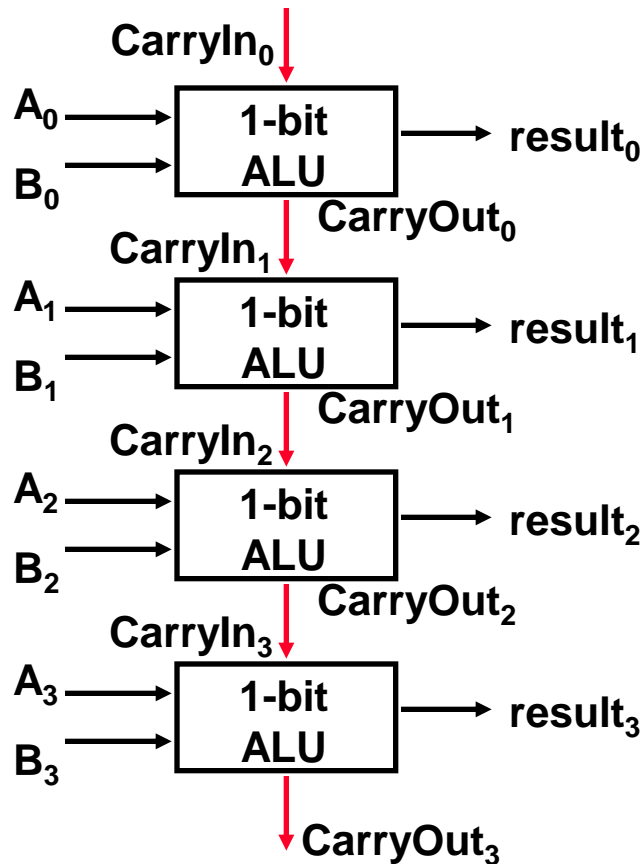
Inputs			Outputs	
A	B	C_{in}	S	C_{out}
0	0	0	0	0
0	0	1	1	0
0	1	0	1	0
0	1	1	0	1
1	0	0	1	0
1	0	1	0	1
1	1	0	0	1
1	1	1	1	1

❑ $S = C_{in} \oplus (A \oplus B)$

❑ $C_{out} = AB + BC_{in} + AC_{in}$

Adder implementation

❑ N-bit ripple-carry adder



Performance depends on data length

➔ Performance is low

Making addition faster: infinite hardware

❑ Parallelize the adder with the cost of hardware

❑ Given the addition:

$$a_{n-1}a_{n-2} \dots a_1a_0 + b_{n-1}b_{n-2} \dots b_1b_0$$

❑ Let c_i is the carry at bit i

$$c_2 = (b_1 \cdot c_1) + (a_1 \cdot c_1) + (a_1 \cdot b_1)$$

$$c_1 = (b_0 \cdot c_0) + (a_0 \cdot c_0) + (a_0 \cdot b_0)$$

Find c_2 from a_0, b_0, a_1, b_1 ?

Making addition faster: Carry Lookahead

- ❑ Video demo:

<https://www.youtube.com/watch?v=yj6wo5SCObY>

- ❑ Approach

- ❑ Make hardwired 4 bit adder → fast and simple enough
- ❑ Develop a carry lookahead unit to calculate the carry bit before finishing the addition

- ❑ At bit i

$$\begin{aligned} c_{i+1} &= (b_i \cdot c_i) + (a_i \cdot c_i) + (a_i \cdot b_i) \\ &= (a_i \cdot b_i) + (a_i + b_i) \cdot c_i \end{aligned}$$

$$g_i = a_i \cdot b_i$$

$$p_i = a_i + b_i$$

- ❑ Denote

$$c_{i+1} = g_i + p_i \cdot c_i$$

- ❑ Then

Carry lookahead

❑ With 4-bit adder

$$c1 = g0 + (p0 \cdot c0)$$

$$c2 = g1 + (p1 \cdot g0) + (p1 \cdot p0 \cdot c0)$$


$$c3 = g2 + (p2 \cdot g1) + (p2 \cdot p1 \cdot g0) + (p2 \cdot p1 \cdot p0 \cdot c0)$$

$$c4 = g3 + (p3 \cdot g2) + (p3 \cdot p2 \cdot g1) + (p3 \cdot p2 \cdot p1 \cdot g0) \\ + (p3 \cdot p2 \cdot p1 \cdot p0 \cdot c0)$$

- ➔ All carry bits can be calculated after 3 gate delay
- ➔ All result bits can be calculated after maximum of 4 gate delay
- ➔ How to implement bigger adder?

Carry lookahead

- For 16-bit adder → fast C1, C2, C3, C4 is needed



Carry lookahead

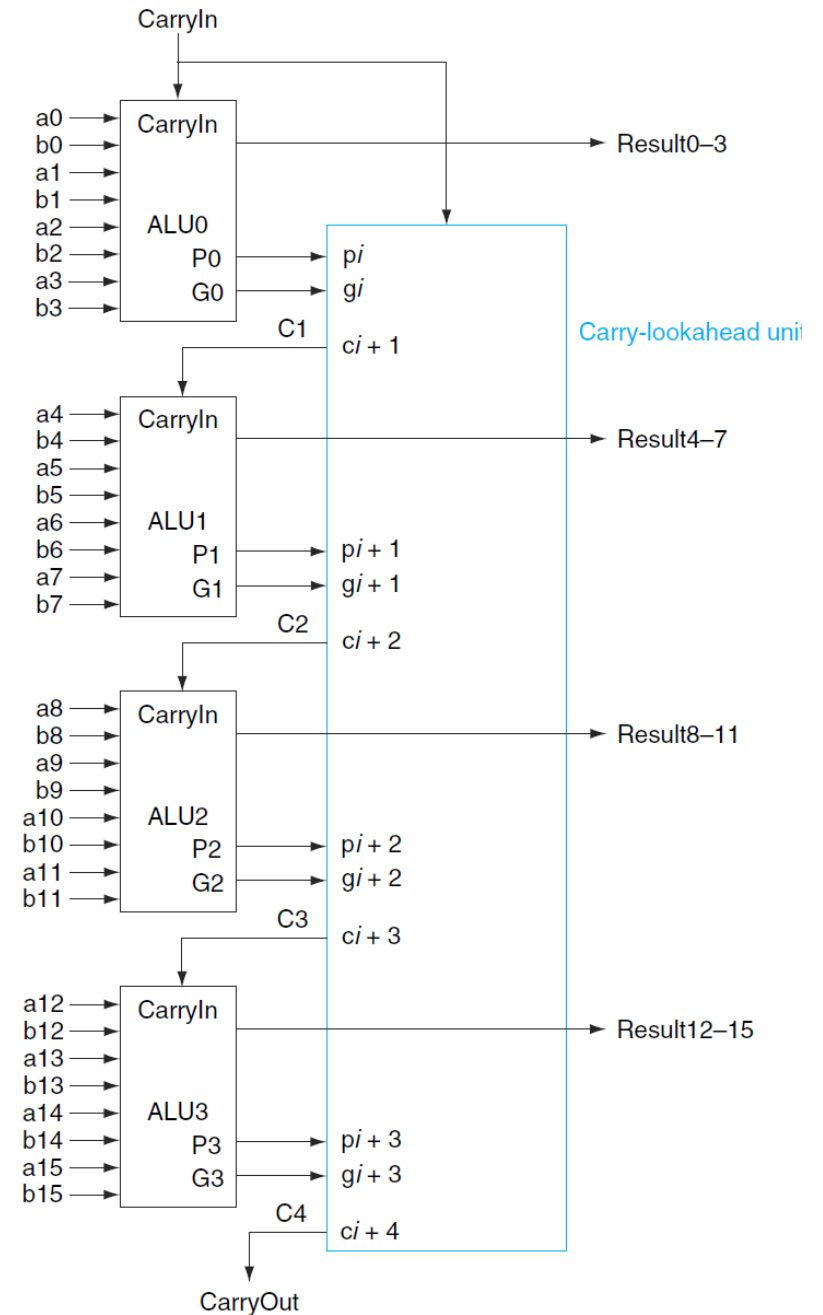
□ Denote

$$\begin{aligned}P_0 &= p_3 \cdot p_2 \cdot p_1 \cdot p_0 & G_0 &= g_3 + (p_3 \cdot g_2) + (p_3 \cdot p_2 \cdot g_1) + (p_3 \cdot p_2 \cdot p_1 \cdot g_0) \\P_1 &= p_7 \cdot p_6 \cdot p_5 \cdot p_4 & G_1 &= g_7 + (p_7 \cdot g_6) + (p_7 \cdot p_6 \cdot g_5) + (p_7 \cdot p_6 \cdot p_5 \cdot g_4) \\P_2 &= p_{11} \cdot p_{10} \cdot p_9 \cdot p_8 & G_2 &= g_{11} + (p_{11} \cdot g_{10}) + (p_{11} \cdot p_{10} \cdot g_9) + (p_{11} \cdot p_{10} \cdot p_9 \cdot g_8) \\P_3 &= p_{15} \cdot p_{14} \cdot p_{13} \cdot p_{12} & G_3 &= g_{15} + (p_{15} \cdot g_{14}) + (p_{15} \cdot p_{14} \cdot g_{13}) + (p_{15} \cdot p_{14} \cdot p_{13} \cdot g_{12})\end{aligned}$$

□ Then big-carry bits can be calculated fast

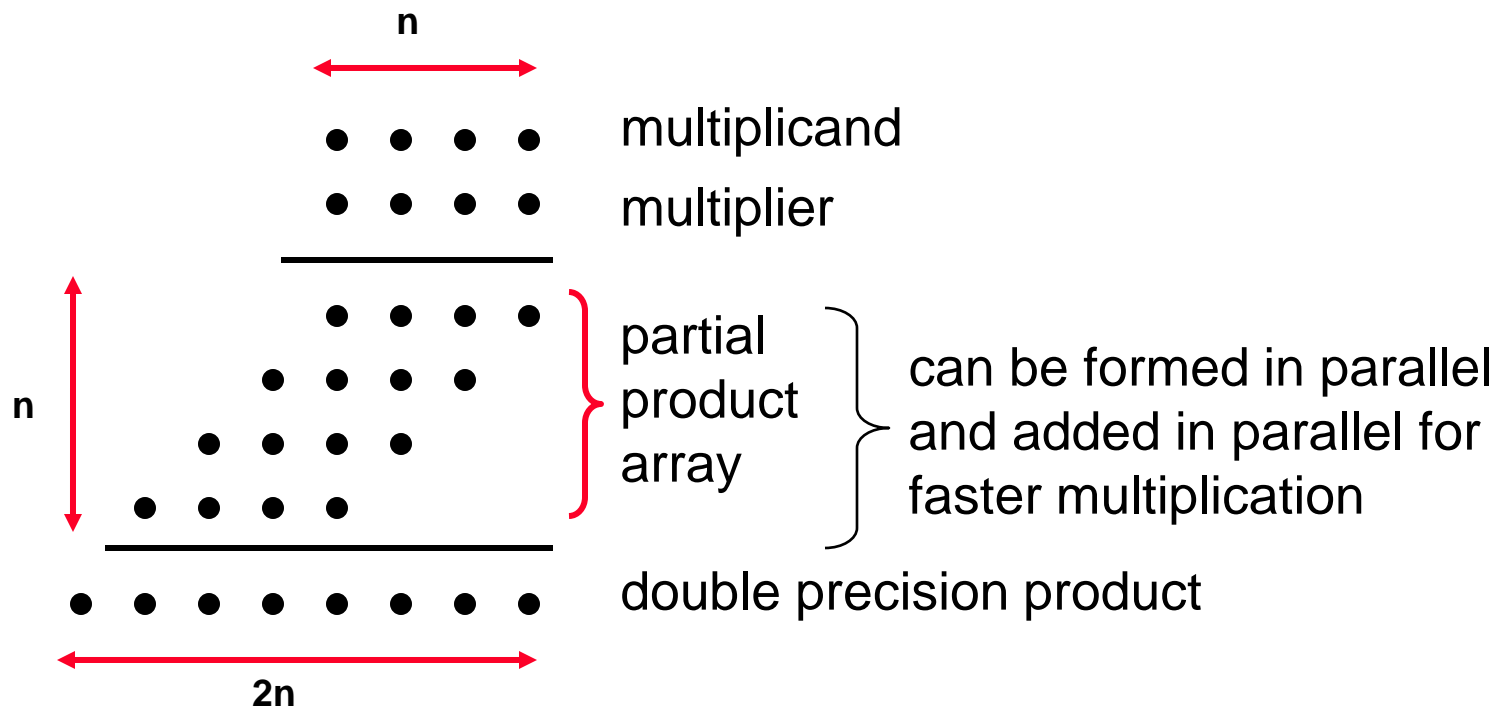
$$\begin{aligned}C_1 &= G_0 + (P_0 \cdot c_0) \\C_2 &= G_1 + (P_1 \cdot G_0) + (P_1 \cdot P_0 \cdot c_0) \\C_3 &= G_2 + (P_2 \cdot G_1) + (P_2 \cdot P_1 \cdot G_0) + (P_2 \cdot P_1 \cdot P_0 \cdot c_0) \\C_4 &= G_3 + (P_3 \cdot G_2) + (P_3 \cdot P_2 \cdot G_1) + (P_3 \cdot P_2 \cdot P_1 \cdot G_0) \\&\quad + (P_3 \cdot P_2 \cdot P_1 \cdot P_0 \cdot c_0)\end{aligned}$$

16-bit Adder



Multiply

- ❑ Binary multiplication is just a *bunch* of right shifts and adds



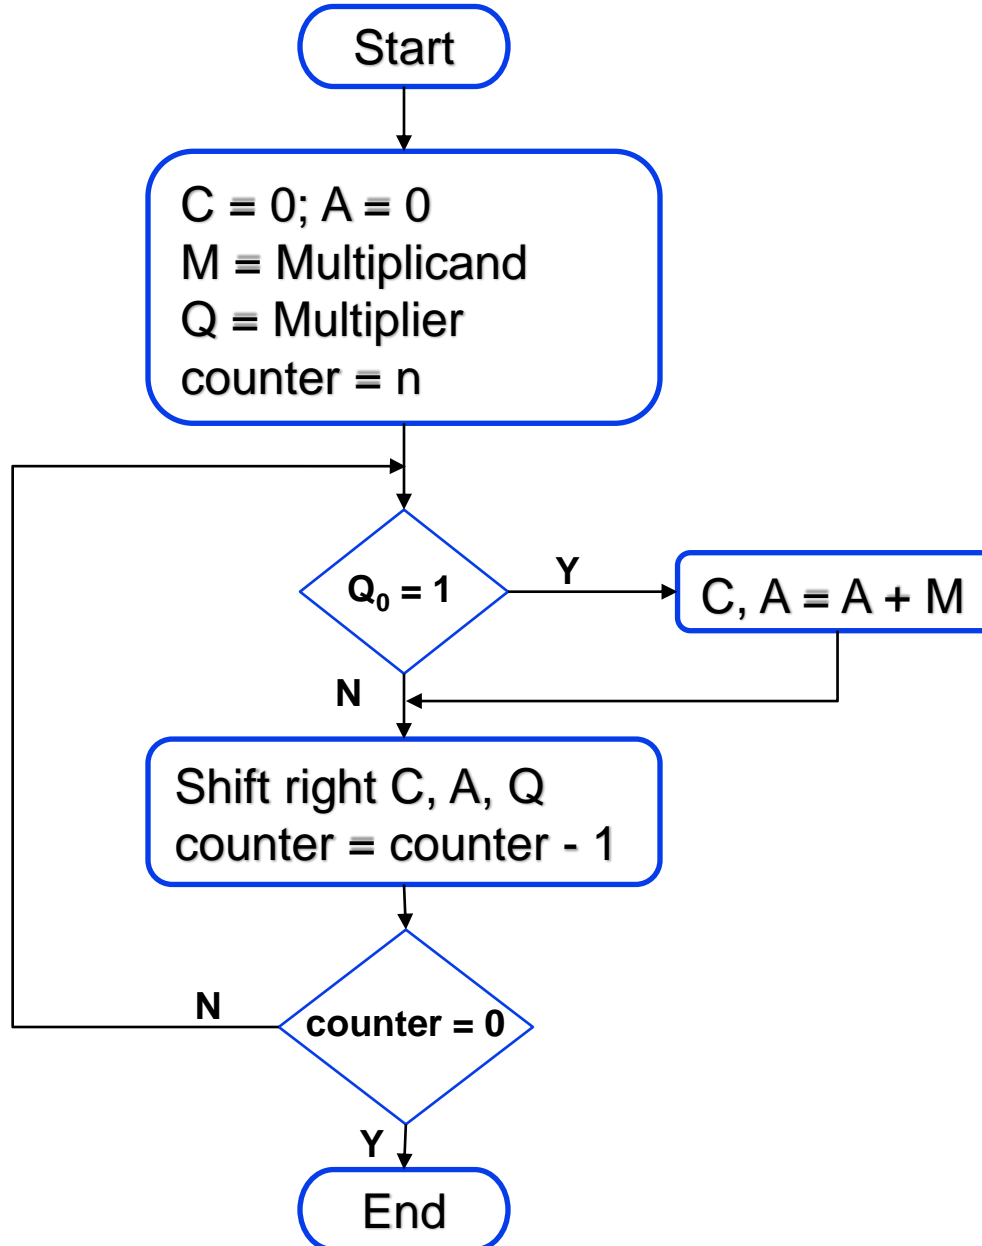
n -bit multiplicand and multiplier \rightarrow $2n$ -bit product

Example

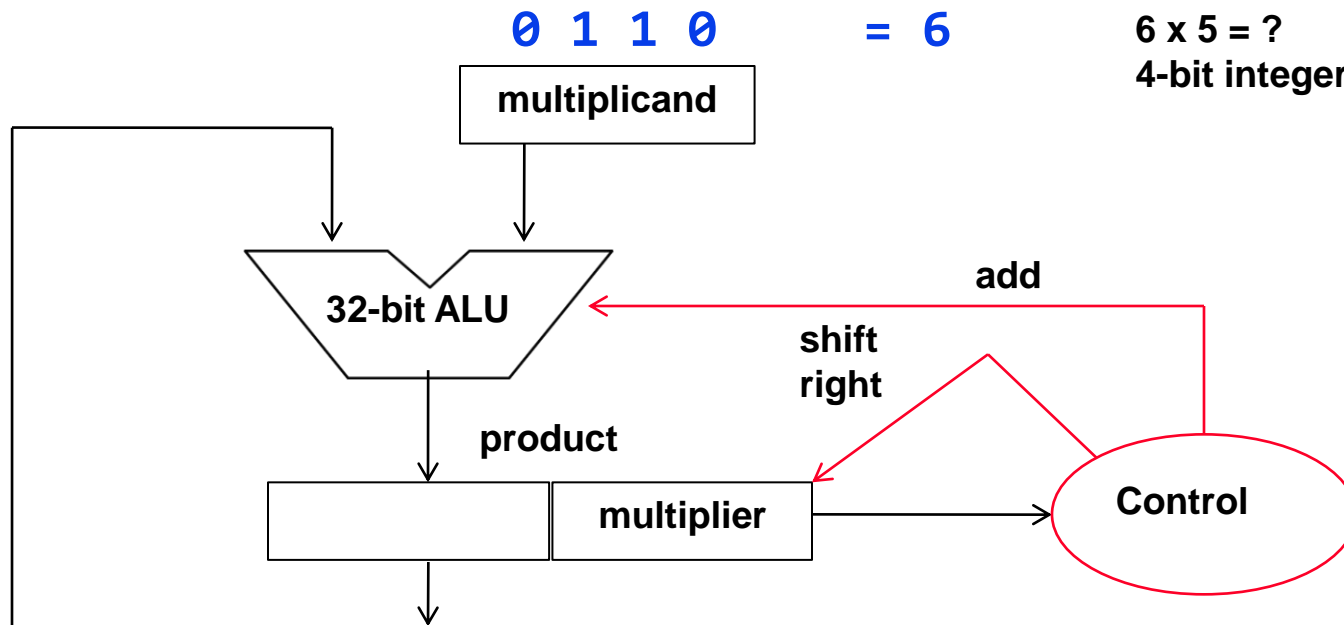
$$\begin{array}{r} \text{Multiplicand} \qquad 1000_{\text{ten}} \\ \text{Multiplier} \quad \times \quad 1001_{\text{ten}} \\ \hline \qquad 1000 \\ \qquad 0000 \\ \qquad 0000 \\ \qquad 1000 \\ \hline \text{Product} \qquad 1001000_{\text{ten}} \end{array}$$

How to do this in hardware?

Add and Right Shift Multiplier



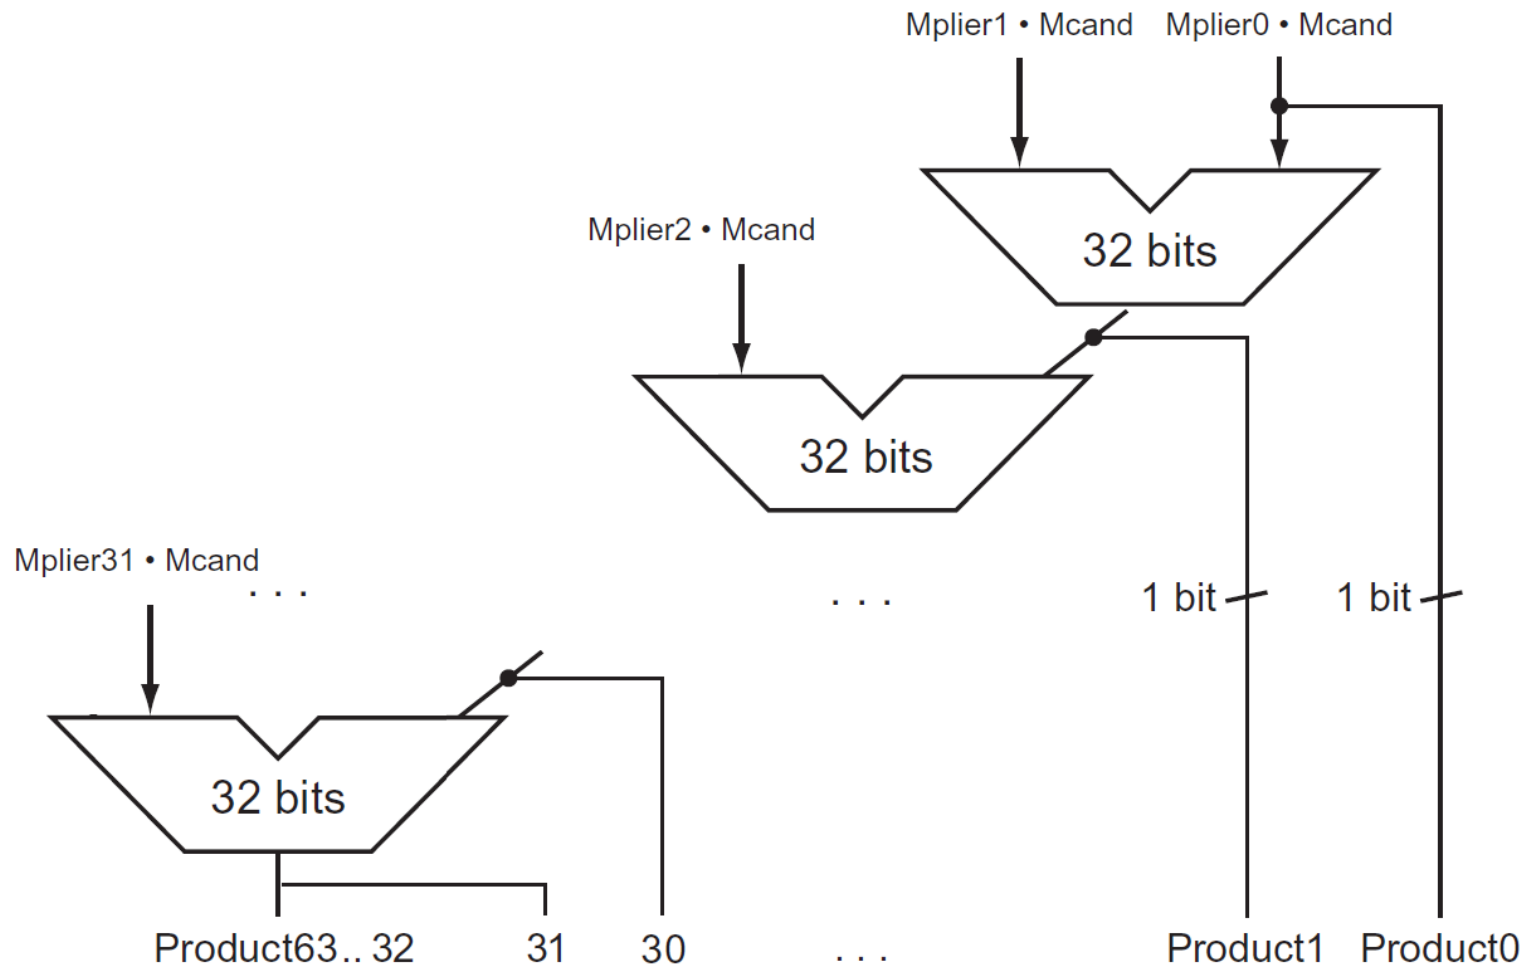
Add and Right Shift Multiplier Hardware



	0 0 0 0	0 1 0 1	= 5
add	0 1 1 0	0 1 0 1	LSB=1 → add multiplicand
	0 0 1 1 →	0 0 1 0	shift right
add	0 0 1 1	0 0 1 0	LSB=0 → no change
	0 0 0 1 →	1 0 0 1	shift right
add	0 1 1 1	1 0 0 1	LSB=1 → add multiplicand
	0 0 1 1 →	1 1 0 0	shift right
add	0 0 1 1	1 1 0 0	LSB=0 → no change
	0 0 0 1 →	1 1 1 0	shift right
		= 30	


Fast multiplier – Design for Moore

❑ Why is this fast?



Fast multiplier – Design for Moore

- ❑ How fast is this?
- ❑ Note: the size of addition circuits



RISC-V Multiply Instruction (RV32M extension)

❑ Multiply instructions: `mul`, `mulh` `mulhu`, `mulhsu`

`mul t1, s0, s1` #set t1 to lower 32 bits
of `s0 * s1`

`mulh t1, s0, s1` #set t1 to upper 32 bits
of `s0 * s1`


`mulhu t1, s0, s1` #set t1 to upper 32 bits of
`s0 * s1` (unsigned multiplication)

`mulhsu t1, s0, s1` #set t1 to upper 32 bits
of `s0 * s1`, where `s0` is signed and `s1` is
unsigned

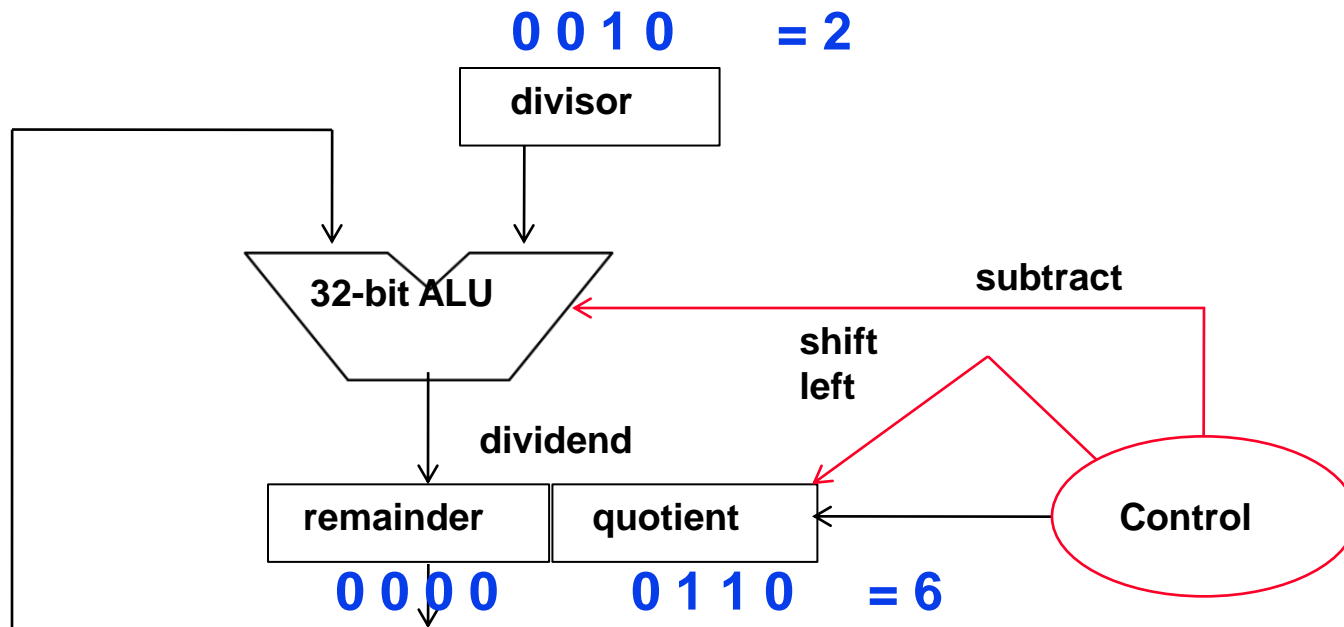
Division

- Division is just a *bunch* of quotient digit guesses and left shifts and subtracts

$$\text{dividend} = \text{quotient} \times \text{divisor} + \text{remainder}$$



Left Shift and Subtract Division Hardware



	0 0 0 0	←	1 1 0 0	
sub	1 1 1 0		1 1 0 0	0
	0 0 0 0		1 1 0 0	
	0 0 0 1	←	1 0 0 0	
sub	1 1 1 1		1 0 0 0	0
	0 0 0 1		1 0 0 0	
	0 0 1 1	←	0 0 0 0	
sub	0 0 0 1		0 0 0 0	1
	0 0 1 0	←	0 0 1 0	
sub	0 0 0 0		0 0 1 0	1

rem < 0, so quotient bit = 0 and restore remainder

rem < 0, so quotient bit = 0 and restore remainder

rem ≥ 0, so quotient bit = 1

rem ≥ 0, so quotient bit = 1

all bits of dividend are processed

Result: remainder 0, quotient 3

RISC-V Divide Instruction (RV32M extension)

❑ Instructions:

`div t1, t2, t3 #t1 = t2/t3 (signed division)`

`divu t1, t2, t3 #t1 = t2/t3 (unsigned division)`

`rem t1, t2, t3 #t1 = remainder of t2/t3`

`remu t1, t2, t3 #t1 = remainder of t2/t3
(unsigned)`

- ❑ As with multiply, divide ignores overflow so software must determine if the quotient is too large.
- ❑ Software must also check the divisor to avoid division by 0.

Signed integer multiplication and division

- ❑ Reuse unsigned multiplication then fix product sign later
- ❑ Multiplication
 - ❑ Multiplicand and multiplier are of the same sign: keep product
 - ❑ Multiplicand and multiplier are of different sign: negate product
- ❑ Division:
 - ❑ Dividend and divisor of the same sign:
 - Keep quotient
 - Keep/negate remainder so it is of the same sign with dividend
 - ❑ Dividend and divisor of different sign:
 - Negate quotient
 - Keep/negate remainder so it is of the same sign with dividend

Example

- ❑ Write a RISC-V program
 - ❑ Reads 2 integers a and b from console
 - ❑ Print out the two values: (a / b) and (a % b) to console

Exercise

- ❑ Write a program that
 - ❑ Reads two integers a and b from console.
 - ❑ Find and print out the greatest common divisor of a and b .
 - ❑ Find and print out the least common multiplier of a and b .

Representing Big (and Small) Numbers

❑ Encoding non-integer value?

- [illegible]

- PI number

$$\text{PI} = 3.14159\dots$$

❑ Problem: how to represent the above numbers?

➔ We need reals or floating-point numbers!

➔ Floating point numbers in decimal:

→ 1000

→ 1×10^3

→ 0.1×10^4

Floating point number

- ❑ In decimal system

$$2013.1228 = 201.31228 * 10$$

$$= 20.131228 * 10^2$$

$$= 2.0131228 * 10^3$$

$$= 20131228 * 10^{-4}$$

- ❑ What is the “standard” form?

$$2.0131228 * 10^3 = \underline{2.0131228} \underline{E+03}$$

mantissa

exponent

- ❑ In binary $X = \pm 1.xxxxx * 2^{yyyy}$


- ❑ ***Sign, mantissa, and exponent need to be represented***

Floating point number

- ❑ Defined by the IEEE 754-1985 standard
 - ❑ Single precision: 32 bit
 - ❑ Double precision: 64 bit
 - ❑ Correspond to float and double in C
- ❑ Single precision floating point representation

$$(-1)^{\text{sign}} \times 1.F \times 2^{E-\text{bias}}$$

- ❑ Fit everything in 32 bits
- ❑ Bias = 127 (with single precision)



Examples

❑ Ex1: convert X into decimal value

$X = 1\textcolor{red}{100}\textcolor{red}{0001}\textcolor{red}{0}101\ 0110\ 0000\ 0000\ 0000\ 0000$

sign = 1 \rightarrow X is negative

E = 1000 0010 = 130

F = 10101100...00

$\rightarrow X = (-1)^1 \times 1.101011000..00 \times 2^{130-127}$

$= -1.101011 \times 2^3 = -1101.011$

$= -13.375$

Example

❑ Ex2: find decimal value of X

X = 0011 1111 1000 0000 0000 0000 0000 0000

sign = 0

e = 0111 1111 = 127

m = 000...0000 (23 bit 0)

$X = (-1)^0 \times 1.00...000 \times 2^{127-127} = 1.0$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

Converting X to plain binary

$$9_{10} = 1001_2$$

$$0.6875 \times 2 = 1.375 \quad \rightarrow \text{get bit } 1$$

$$0.375 \times 2 = 0.75 \quad \rightarrow \text{get bit } 0$$

$$0.75 \times 2 = 1.5 \quad \rightarrow \text{get bit } 1$$

$$0.5 \times 2 = 1.0 \quad \rightarrow \text{get bit } 1$$



$$\rightarrow 9.6875_{10} = 1001.1011_2$$

Example

- Ex3: find binary representation of $X = 9.6875$ in IEEE 754 single precision

$$X = 9.6875_{(10)} = 1001.1011_{(2)} = 1.0011011 \times 2^3$$

Then

$$S = 0$$

$$e = 127 + 3 = 130_{(10)} = 1000\ 0010_{(2)}$$

$$m = 001101100\dots00 \text{ (23 bit)}$$

Finally

$$X = 0100\ 0001\ 0001\ 1011\ 0000\ 0000\ 0000\ 0000$$

Examples

□ $1.0_2 \times 2^{-1} =$


□ $100.75_{10} =$

Some special values


- ❑ Smallest+: 0 00000001 1.00000000000000000000000000000000
= $1 \times 2^{1-127}$
- ❑ Zero: 0 00000000 00000000000000000000000000000000
= true 0
- ❑ Largest+: 0 11111110 1.11111111111111111111111111111111
= $(2-2^{-23}) \times 2^{254-127}$

Too large or too small values

- ❑ **Overflow** (floating point) happens when a positive exponent becomes too large to fit in the exponent field
- ❑ **Underflow** (floating point) happens when a negative exponent becomes too large to fit in the exponent field



- ❑ **Double precision: 64 bits**



IEEE 754 FP Standard Encoding

- ❑ Special encodings are used to represent unusual events
 - ❑ \pm infinity for division by zero
 - ❑ NaN (not a number) for invalid operations such as 0/0
 - ❑ True zero is the bit string all zero

Single Precision		Double Precision		Object Represented
E (8)	F (23)	E (11)	F (52)	
0000 0000	0	0000 ... 0000	0	true zero (0)
0000 0000	nonzero	0000 ... 0000	nonzero	\pm denormalized number
0111 1111 to +127,-126	anything	0111 ...1111 to +1023,-1022	anything	\pm floating point number
1111 1111	+ 0	1111 ... 1111	- 0	\pm infinity
1111 1111	nonzero	1111 ... 1111	nonzero	not a number (NaN)

Floating Point Addition

□ Addition (and subtraction)

$$(\pm F1 \times 2^{E1}) + (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- **Step 0: Restore the hidden bit in F1 and in F2**
- **Step 1: Align fractions by right shifting F2 by $E1 - E2$ positions (assuming $E1 \geq E2$) keeping track of (three of) the bits shifted out in G R and S**
- **Step 2: Add the resulting F2 to F1 to form F3**
- **Step 3: Normalize F3 (so it is in the form 1.XXXXXX ...)**
 - If F1 and F2 have the same sign $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment $E3$ (check for overflow)
 - If F1 and F2 have different signs $\rightarrow F3$ may require *many* left shifts each time decrementing $E3$ (check for underflow)
- **Step 4: Round F3 and possibly normalize F3 again**
- **Step 5: Rehide the most significant bit of F3 before storing the result**

Floating Point Addition Example

□ Add

$$(0.5 = 1.0000 \times 2^{-1}) + (-0.4375 = -1.1100 \times 2^{-2})$$

□ **Step 0:**

□ **Step 1:**

□ **Step 2:**

□ **Step 3:**

□ **Step 4:**

□ **Step 5:**

Floating Point Addition Example

□ Add: $0.5 + (-0.4375) = ?$

$$(0.5 = 1.0000 \times 2^{-1}) + (-0.4375 = -1.1100 \times 2^{-2})$$

- **Step 0:** Hidden bits restored in the representation above
- **Step 1:** Shift significand with the smaller exponent (1.1100) right until its exponent matches the larger exponent (so once)
- **Step 2:** Add significands
$$1.0000 + (-0.111) = 1.0000 - 0.111 = 0.001$$
- **Step 3:** Normalize the sum, checking for exponent over/underflow
$$0.001 \times 2^{-1} = 0.010 \times 2^{-2} = \dots = 1.000 \times 2^{-4}$$
- **Step 4:** The sum is already rounded, so we're done
- **Step 5:** Rehide the hidden bit before storing

Floating Point Multiplication

❑ Multiplication

$$(\pm F1 \times 2^{E1}) \times (\pm F2 \times 2^{E2}) = \pm F3 \times 2^{E3}$$

- ❑ **Step 0: Restore the hidden bit in F1 and in F2**
- ❑ **Step 1: Add the two (biased) exponents and subtract the bias from the sum, so $E1 + E2 - 127 = E3$**
also determine the sign of the product (which depends on the sign of the operands (most significant bits))
- ❑ **Step 2: Multiply F1 by F2 to form a double precision F3**
- ❑ **Step 3: Normalize F3 (so it is in the form 1.XXXXXX ...)**
 - Since F1 and F2 come in normalized $\rightarrow F3 \in [1,4) \rightarrow$ 1 bit right shift F3 and increment E3
 - Check for overflow/underflow
- ❑ **Step 4: Round F3 and possibly normalize F3 again**
- ❑ **Step 5: Rehide the most significant bit of F3 before storing the result**

Floating Point Multiplication Example

❑ Multiply

$$(0.5 = 1.0000 \times 2^{-1}) \times (-0.4375 = -1.1100 \times 2^{-2})$$

❑ **Step 0:**

❑ **Step 1:**

❑ **Step 2:**

❑ **Step 3:**

❑ **Step 4:**

❑ **Step 5:**

Floating Point Multiplication Example

❑ Multiply

$$(0.5 = 1.0000 \times 2^{-1}) \times (-0.4375 = -1.1100 \times 2^{-2})$$

❑ **Step 0:** Hidden bits restored in the representation above

❑ **Step 1:** Add the exponents (not in bias would be $-1 + (-2) = -3$ and in bias would be $(-1+127) + (-2+127) - 127 = (-1 -2) + (127+127-127) = -3 + 127 = 124$)

❑ **Step 2:** Multiply the significands

$$1.0000 \times 1.110 = 1.110000$$

❑ **Step 3:** Normalized the product, checking for exp over/underflow

$$1.110000 \times 2^{-3} \text{ is already normalized}$$

❑ **Step 4:** The product is already rounded, so we're done

❑ **Step 5:** Rehide the hidden bit before storing

Support for Accurate Arithmetic

- ❑ Rounding (except for truncation) requires the hardware to include extra F bits during calculations
 - ❑ Guard and Round bit – 2 additional bits to increase accuracy
 - ❑ Sticky bit – used to support **Round to nearest even**; is set to a 1 whenever a 1 bit shifts (right) through it (e.g., when aligning F during addition/subtraction)

F = 1 . xxxxxxxxxxxxxxxxxxxxxxxxx G R S

- ❑ IEEE 754 FP rounding modes
 - ❑ Always round up (toward $+\infty$)
 - ❑ Always round down (toward $-\infty$)
 - ❑ Truncate
 - ❑ **Round to nearest even** (when the Guard || Round || Sticky are 100) – always creates a 0 in the least significant (kept) bit of F

<http://pages.cs.wisc.edu/~markhill/cs354/Fall2008/notes/flpt.apprec.html>

Example

❑ Calculate:

$$0.2 \times 5 = ?$$

$$0.333 \times 3 = ?$$

$$(1.0/3) \times 3 = ?$$

Floating point instructions: RV32F

RV32F / D Floating-Point Extensions

Inst	Name	FMT	Opcode	funct3	funct5	Description (C)
flw	Flt Load Word	*				<code>rd = M[rs1 + imm]</code>
fsw	Flt Store Word	*				<code>M[rs1 + imm] = rs2</code>
fmadd.s	Flt Fused Mul-Add	*				<code>rd = rs1 * rs2 + rs3</code>
fmsub.s	Flt Fused Mul-Sub	*				<code>rd = rs1 * rs2 - rs3</code>
fnmadd.s	Flt Neg Fused Mul-Add	*				<code>rd = -rs1 * rs2 + rs3</code>
fnmsub.s	Flt Neg Fused Mul-Sub	*				<code>rd = -rs1 * rs2 - rs3</code>
fadd.s	Flt Add	*				<code>rd = rs1 + rs2</code>
fsub.s	Flt Sub	*				<code>rd = rs1 - rs2</code>
fmul.s	Flt Mul	*				<code>rd = rs1 * rs2</code>
fdiv.s	Flt Div	*				<code>rd = rs1 / rs2</code>
fsqrt.s	Flt Square Root	*				<code>rd = sqrt(rs1)</code>
fsgnj.s	Flt Sign Injection	*				<code>rd = abs(rs1) * sgn(rs2)</code>
fsgnjn.s	Flt Sign Neg Injection	*				<code>rd = abs(rs1) * -sgn(rs2)</code>
fsgnjx.s	Flt Sign Xor Injection	*				<code>rd = rs1 * sgn(rs2)</code>
fmin.s	Flt Minimum	*				<code>rd = min(rs1, rs2)</code>
fmax.s	Flt Maximum	*				<code>rd = max(rs1, rs2)</code>
fcvt.s.w	Flt Conv from Sign Int	*				<code>rd = (float) rs1</code>
fcvt.s.wu	Flt Conv from Uns Int	*				<code>rd = (float) rs1</code>
fcvt.w.s	Flt Convert to Int	*				<code>rd = (int32_t) rs1</code>
fcvt.wu.s	Flt Convert to Int	*				<code>rd = (uint32_t) rs1</code>
fmv.x.w	Move Float to Int	*				<code>rd = *((int*) &rs1)</code>
fmv.w.x	Move Int to Float	*				<code>rd = *((float*) &rs1)</code>
feq.s	Float Equality	*				<code>rd = (rs1 == rs2) ? 1 : 0</code>
flt.s	Float Less Than	*				<code>rd = (rs1 < rs2) ? 1 : 0</code>
fle.s	Float Less / Equal	*				<code>rd = (rs1 <= rs2) ? 1 : 0</code>
fclass.s	Float Classify	*				<code>rd = 0..9</code>

Exercise

- ❑ Write the corresponding RISC-V assembly program equivalent to the following C code:

```
float x = 0.75;  
printf("%f", x);
```