

25 YEARS ANNIVERSARY
SOICT

HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY



HA NOI UNIVERSITY OF SCIENCE AND TECHNOLOGY
SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

Lesson 6

Storage and Index

Outline

- Overview of database storage structures
- Physical database files
- Database index

Objectives

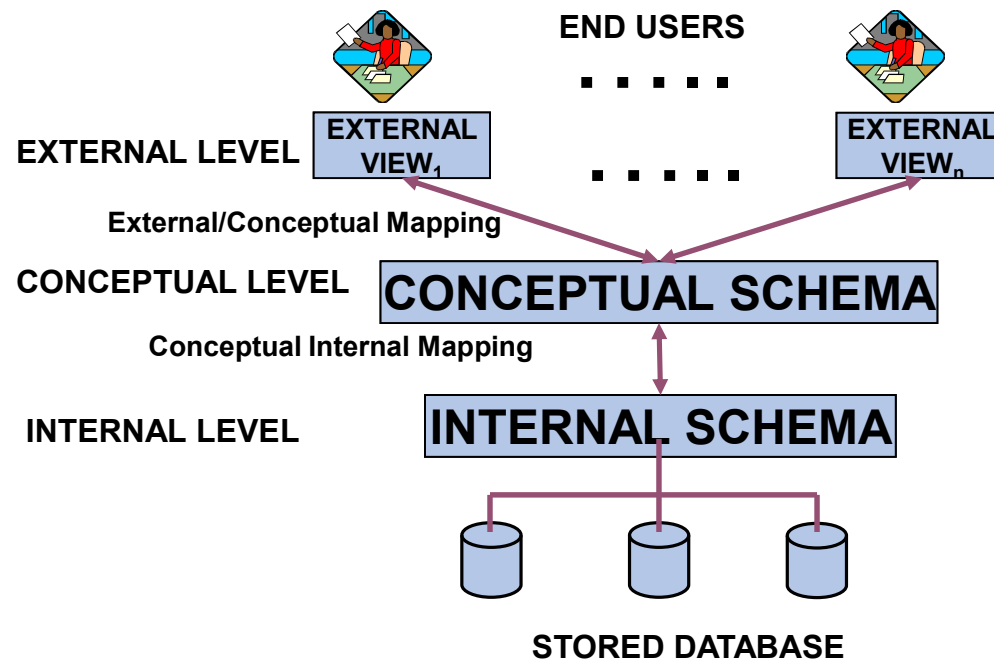
- Upon completion of this lesson, students will be able to:
 - Understand the physical database files
 - Understand the role of database indexes

Keywords

Heap file	Files of Unordered Records
Ordered file	Physically order the records of a file on disk based on the values of one of their fields (key field)
Index	A data structure that improves the speed of data retrieval operations
B-tree	A self-balancing tree data structure that keeps data sorted

1. Overview of database storage structures

3-tier Schema Model (ANSI-SPARC Architecture)



1. Overview of database storage structures

How does Mariadb store data

```
MariaDB [(none)]> SHOW VARIABLES LIKE 'datadir';
```

Variable_name	Value
datadir	/var/lib/mysql/

```
MariaDB [student_management]> show tables;
```

Tables_in_student_management
class
enrolled
faculty
student

the .frm table file stores the table's format
the .ibd file stores the table's data

```
:/var/lib/mysql/student_management# ls -la
```

ql mysql	4096	Mar 12 02:05	.
ql mysql	4096	May 5 06:06	
ql mysql	1547	Mar 12 02:05	class.frm
ql mysql	114688	Mar 12 02:21	class.ibd
ql mysql	65	Mar 12 01:59	db.opt
ql mysql	1466	Mar 12 02:03	enrolled.frm
ql mysql	114688	Mar 12 02:18	enrolled.ibd
ql mysql	1005	Mar 12 02:04	faculty.frm
ql mysql	98304	Mar 12 02:16	faculty.ibd
ql mysql	1101	Mar 12 02:00	student.frm
ql mysql	98304	Mar 12 02:23	student.ibd



SCHOOL OF INFORMATION AND COMMUNICATION TECHNOLOGY

1. Overview of database storage structures

How does Mariadb store data

- the .frm file stores the table's format

```
MariaDB [student_management]> describe student;
```

Field	Type	Null	Key	Default	Extra
snum	int(11)	NO	PRI	NULL	
sname	varchar(40)	YES		NULL	
major	varchar(30)	YES		NULL	
level	varchar(10)	YES		NULL	
age	int(11)	YES		NULL	

```
root@285e07e9458f:/var/lib/mysql/student_management# cat student.frm
?
```

```
VM?\!  ?s?$??%?解 B??
??PRIMARY??InnoDB??f\P
(/?
```

```
N?
```

```
?snum?sname?major?level?age?root@285e07e9458f:/var/lib/mysql/student
```


1. Overview of database storage structures

How does Mariadb store data

- the .ibd file stores the table's data

```
MariaDB [student_management]> select * from student;
```

snum	sname	major	level	age
1	Nguyen Van A	CS	JR	18
2	Nguyen Viet Cuong	History	JR	19
3	Nguyen Hong Ngoc	CS	JR	19
4	Mark Juke	History	JR	20
5	Elon Mulk	CS	JR	20
6	Donal Trump	CS	JR	20
7	Obama	CS	JR	20
8	Tan Dung	History	SR	30

```
[root@285e07e9458f:/var/lib/mysql/student_management# cat student.ibd
???]&!????????????????????????????????????????????????????????????&[?]Y?&??Y?&???j?&[?]??
????????????????????????????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????????????????????????????
????????????????i????????????????????????????????????????????????????????????????
????????????????????????????????????????????????????????????????????????????????
9infimum
supremum
.??8?WNguyen Van ACSJR?8?:?cNguyen Viet CuongHistoryJR? 2?@??Nguyen H
ong NgocCSJR? (0?1?1Mark JukeHistoryJR? 0+?0??Elon MulkCSJR?
8-?Q?kDonal Trum
pCSJR?'?W??ObamaCSJRH????Tan DungHistorySR?pc??Q?'??root@285e07e9458f:/var/lib/mys
```

2. Physical database files

Motivation

Magnetic disks as data storage

Primary file organizations

2.1. Motivation

- Databases typically store large amounts of data persistently on disks:
 - Databases are too large to fit entirely in main memory.
 - Disk - nonvolatile storage vs. Main memory - volatile storage
 - The cost of storage per unit is much cheaper

2.2. Magnetic disks as data storage

- A disk is a random access addressable device.
- Transfer of data between main memory and disk takes place in units of disk blocks.
- Typical disk block sizes: 4KB – 8KB.
- Disk I/O (read/write from disk to main memory) overhead is the key factor of database performance optimization.

2.2.1. Physical database design

- The process of physical database design involves choosing the **particular data organization techniques that best suit the given application requirements** (on SELECT, INSERT, UPDATE, DELETE).
- The data stored on disk is organized as files of records:
 - **Primary file organizations**: determine how the file records are physically placed on the disk, and hence how the records can be accessed.
 - **Secondary organization** or auxiliary access structure allows efficient access to file records based on alternate fields.

2.2.2. Placing File Records on Disk

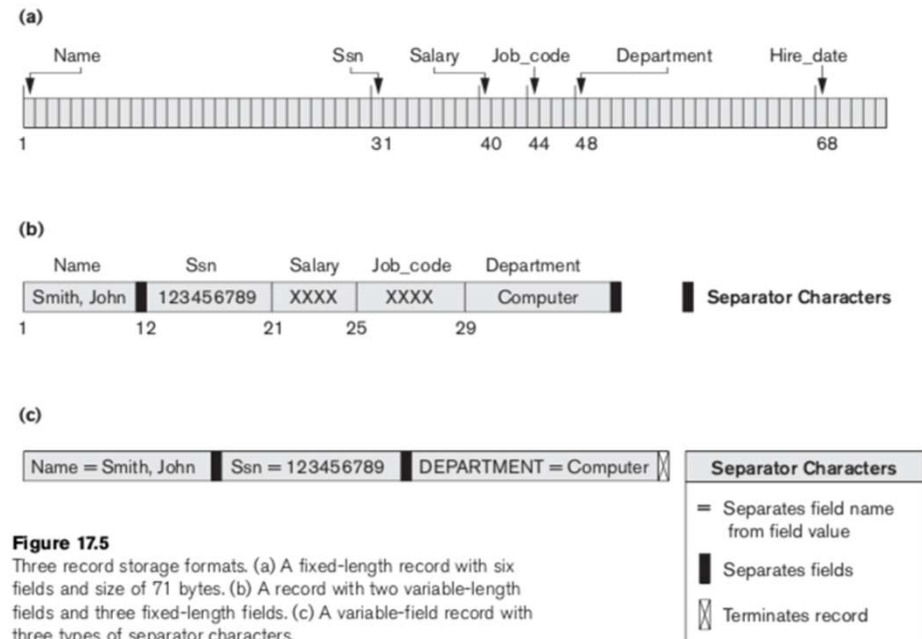


Figure 17.5

Three record storage formats. (a) A fixed-length record with six fields and size of 71 bytes. (b) A record with two variable-length fields and three fixed-length fields. (c) A variable-field record with three types of separator characters.

© Elmasri, Ramez. *Fundamentals of database systems*. Pearson Education India, 2008

2.3. Primary file organizations

- Files of Unordered Records (Heap Files)
- Files of Ordered Records (Sorted Files)
- Hashing Techniques

2.3. Primary file organizations

- Files of Unordered Records (Heap Files)
 - Records are placed in the file in the order in which they are inserted
 - **INSERT**: Inserting a new record is very efficient
 - New records are inserted at the end of the file
 - **UPDATE/SELECT**: Searching for a record on any search condition is not efficient – linear search
 - **DELETE**: leaves unused space in the disk block
 - require periodic reorganization

2.3. Primary file organizations

- Files of Ordered Records (Sorted Files)
 - Physically order the records of a file on disk based on the values of one of their fields (key field)
 - **SELECT:** binary search (very fast)
 - **INSERT/DELETE/UPDATE:** more expensive

2.3. Primary file organizations


- Hash files
 - The address of the disk block in which the record is stored is the result of applying a hash function to the value of a particular field (hash field) of the record.
 - Very fast access to records for search on equality condition on the hash field.

3. Database indexes

1. What is database index?
2. Index data structures
3. B+tree
4. Spare vs. Dense index
5. Clustered vs. Non-clustered index
6. Index creation in SQL

3.1. What is database index?

- Auxiliary access structure (commonly index) allows efficient access to file records based on alternate fields




3.2. Index data structures

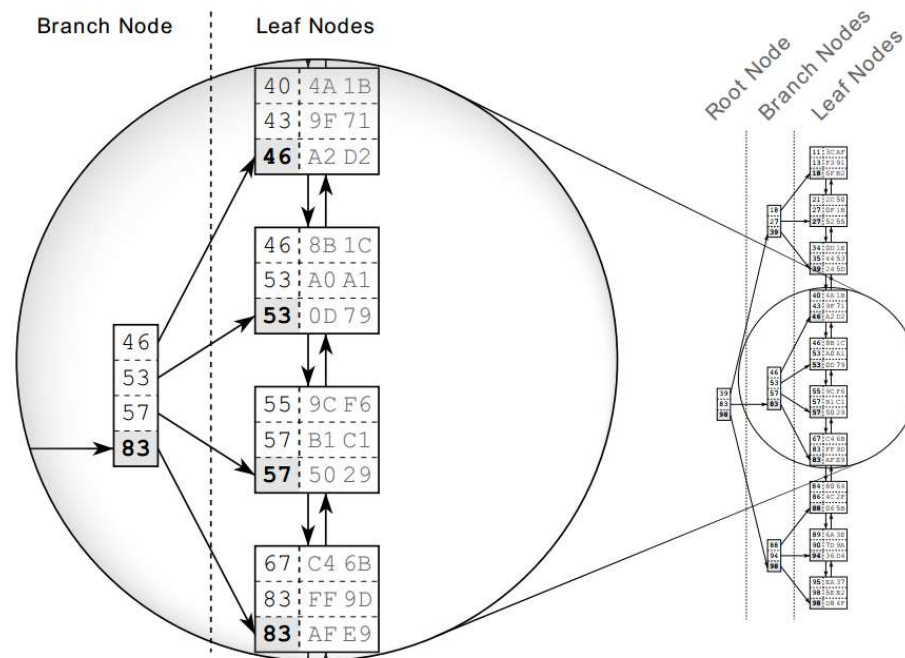
- Indexes can be implemented with different data structures.
 - B+-tree index
 - hash index
 - bitmap index (briefly)
 - dynamic hash indexes: number of buckets modified dynamically
 - R-tree: index for special data (points, lines, shapes)
 - quadtree: recursively partition a 2D plane into four quadrants
 - octree: quadtree version for three dimensional data
 - main memory indexes: T-tree, binary search tree

3.3. B+Tree

- Balanced tree of key-pointer pairs
- Keys are sorted by value
- Nodes are at least half full
- Access records for key: traverse tree from root to leaf




3.3.1. Example: B+ tree



© Gulutzan, Peter, and Trudy Pelzer. *SQL Performance Tuning*. Addison-Wesley Professional, 2003.


3.4. Spare vs. Dense index

- Sparse index
 - pointers to disk pages
 - at most one pointer per disk page
 - usually much less pointers than records
- Dense index
 - pointers to individual records
 - one key per record
 - usually more keys than sparse index optimization: store repeating keys only once, followed by pointers




3.5. Clustered vs. Non-Clustered

- Clustered index on attribute X
 - This index controls the placement of records on disk
 - only one clustering index per table
 - dense or sparse
- Non-clustered index on attribute X
 - no constraint on table organization
 - Can have more than one index per table
 - always dense



3.5.1. Example: Non-clustered index



© Gulutzan, Peter, and Trudy Pelzer. *SQL Performance Tuning*. Addison-Wesley Professional, 2003.

3.6. Creating Index

- CREATE [UNIQUE|FULLTEXT|SPATIAL] INDEX index_name [index_type] ON tbl_name (index_col_name,...) [index_option] [algorithm_option | lock_option] ...
- index_type: USING {BTREE | HASH}

Remark

- Databases typically store data persistently on disks
 - Files of unordered records (Heap files)
 - Files of ordered records (Sorted files)
 - Hash files
- Index allows efficient access to file records based on “indexed” fields

Quiz 1.

Quiz Number	1	Quiz Type	OX	Example Select
Question	Does heap files support INSERT query efficiently?			
Example	A. Yes B. No			
Answer	A			
Feedback	New records are appended to the end of the head file			

Quiz 2.

Quiz Number	2	Quiz Type	OX	Example Select
Question	Are ordered files better for heavy Insert operation?			
Example	A. Yes B. No			
Answer	B			
Feedback	Insertion to ordered files requires reorganizing w.r.t. new records			

Summary

- Overview of database storage structures
 - 3-tier Schema Model (ANSI-SPARC Architecture)
 - How Mariadb stores data
- Physical database file structures
 - Motivation
 - Magnetic disks as data storage
 - Primary file organizations
- Database index
 - What is database indexes?
 - Index data structures
 - B+tree
 - Spare vs. Dense index
 - Clustered vs. Non-clustered index
 - Index creation in SQL

Next lesson: Query processing

- Hector Garcia-Molina, Jeffrey D. Ullman, Jennifer Widom. Database Systems: The Complete Book. Pearson Prentice Hall. the 2nd edition. 2008: Chapter 7
- Nguyen Kim Anh, Nguyên lý các hệ cơ sở dữ liệu, NXB Giáo dục. 2004: Chương 7