

[Open in app](#)[Follow](#)

579K Followers



Creating and training a U-Net model with PyTorch for 2D & 3D semantic segmentation: Model building [2/4]

A guide to semantic segmentation with PyTorch and the U-Net



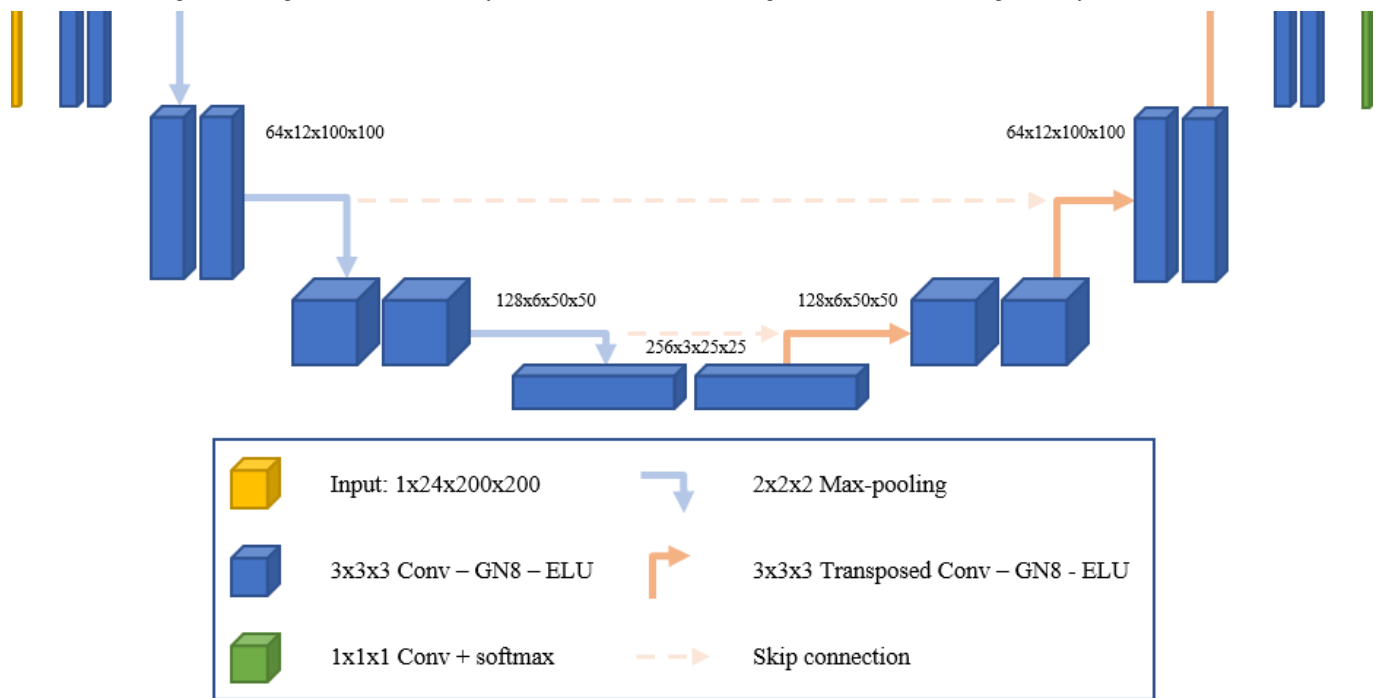
Johannes Schmidt · Dec 2, 2020 · 6 min read

In the [previous chapter](#) we built a dataloader that picks up our images and performs some transformations and augmentations so that they can be fed in batches to a neural network like the U-Net. In this part, we focus on building a U-Net from scratch with the PyTorch library. The goal is to implement the U-Net in such a way, that important model configurations such as the activation function or the depth can be passed as arguments when creating the model.

About the U-Net

The U-Net is a convolutional neural network architecture that is designed for fast and precise segmentation of images. It has performed extremely well in several challenges and to this day, it is one of the most popular end-to-end architectures in the field of semantic segmentation.





The UNet — Image by Johannes Schmidt — Based on <https://arxiv.org/abs/1505.04597>

We can split the network into two parts: The encoder path (backbone) and the decoder path. The encoder captures features at different scales of the images by using a traditional stack of convolutional and max pooling layers. Concretely speaking, a block in the encoder consists of the repeated use of two convolutional layers ($k=3$, $s=1$), each followed by a non-linearity layer, and a max-pooling layer ($k=2$, $s=2$). For every convolution block and its associated max pooling operation, the number of feature maps is doubled to ensure that the network can learn the complex structures effectively.

The decoder path is a symmetric expanding counterpart that uses transposed convolutions. This type of convolutional layer is an up-sampling method with trainable parameters and performs the reverse of (down)pooling layers such as the max pool. Similar to the encoder, each convolution block is followed by such an up-convolutional layer. The number of feature maps is halved in every block. Because recreating a segmentation mask from a small feature map is a rather difficult task for the network, the output after every up-convolutional layer is appended by the feature maps of the corresponding encoder block. The feature maps of the encoder layer are cropped if the dimensions exceed the one of the corresponding decoder layers.

In the end, the output passes another convolution layer ($k=1$, $s=1$) with the number of feature maps being equal to the number of defined labels. The result is a u-shaped

convolutional network that offers an elegant solution for good localization and use of context. Let's take a look at the code.

The code

This code is based on

<https://github.com/ELEKTRONN/elektronn3/blob/master/elektronn3/models/unet.py>

© 2017 Martin Drawitsch, released under MIT License, which implements a configurable (2D/3D) U-Net with user-defined network depth and a few other improvements of the original architecture. They themselves actually used the 2D code from Jackson Huang <https://github.com/jaxony/unet-pytorch>.

Here is a simplified version of the code — saved in a file `unet.py` :

```

1  from torch import nn
2  import torch
3
4
5  @torch.jit.script
6  def autocrop(encoder_layer: torch.Tensor, decoder_layer: torch.Tensor):
7      """
8      Center-crops the encoder_layer to the size of the decoder_layer,
9      so that merging (concatenation) between levels/blocks is possible.
10     This is only necessary for input sizes != 2**n for 'same' padding and always requires
11     """
12     if encoder_layer.shape[2:] != decoder_layer.shape[2:]:
13         ds = encoder_layer.shape[2:]
14         es = decoder_layer.shape[2:]
15         assert ds[0] >= es[0]
16         assert ds[1] >= es[1]
17         if encoder_layer.dim() == 4: # 2D
18             encoder_layer = encoder_layer[
19                 :,
20                 :,
21                 ((ds[0] - es[0]) // 2):((ds[0] + es[0]) // 2),
22                 ((ds[1] - es[1]) // 2):((ds[1] + es[1]) // 2)
23             ]
24         elif encoder_layer.dim() == 5: # 3D
25             assert ds[2] >= es[2]
26             encoder_layer = encoder_layer[
27                 :,
28                 :,
29                 ((ds[0] - es[0]) // 2):((ds[0] + es[0]) // 2),

```

```

29         ((ds[0] - es[0]) // 2, (ds[0] + es[0]) // 2),
30         ((ds[1] - es[1]) // 2, (ds[1] + es[1]) // 2),
31         ((ds[2] - es[2]) // 2, (ds[2] + es[2]) // 2),
32     ]
33     return encoder_layer, decoder_layer
34
35
36 def conv_layer(dim: int):
37     if dim == 3:
38         return nn.Conv3d
39     elif dim == 2:
40         return nn.Conv2d
41
42
43 def get_conv_layer(in_channels: int,
44                   out_channels: int,
45                   kernel_size: int = 3,
46                   stride: int = 1,
47                   padding: int = 1,
48                   bias: bool = True,
49                   dim: int = 2):
50     return conv_layer(dim)(in_channels, out_channels, kernel_size=kernel_size, stride=stride,
51                           padding=padding, bias=bias)
52
53
54 def conv_transpose_layer(dim: int):
55     if dim == 3:
56         return nn.ConvTranspose3d
57     elif dim == 2:
58         return nn.ConvTranspose2d
59
60
61 def get_up_layer(in_channels: int,
62                 out_channels: int,
63                 kernel_size: int = 2,
64                 stride: int = 2,
65                 dim: int = 3,
66                 up_mode: str = 'transposed',
67                 ):
68     if up_mode == 'transposed':
69         return conv_transpose_layer(dim)(in_channels, out_channels, kernel_size=kernel_size,
70                                         stride=stride)
71     else:
72         return nn.Upsample(scale_factor=2.0, mode=up_mode)
73

```

```

74 def maxpool_layer(dim: int):
75     if dim == 3:
76         return nn.MaxPool3d
77     elif dim == 2:
78         return nn.MaxPool2d
79
80
81 def get_maxpool_layer(kernel_size: int = 2,
82                       stride: int = 2,
83                       padding: int = 0,
84                       dim: int = 2):
85     return maxpool_layer(dim=dim)(kernel_size=kernel_size, stride=stride, padding=padding)
86
87
88 def get_activation(activation: str):
89     if activation == 'relu':
90         return nn.ReLU()
91     elif activation == 'leaky':
92         return nn.LeakyReLU(negative_slope=0.1)
93     elif activation == 'elu':
94         return nn.ELU()
95
96
97 def get_normalization(normalization: str,
98                      num_channels: int,
99                      dim: int):
100     if normalization == 'batch':
101         if dim == 3:
102             return nn.BatchNorm3d(num_channels)
103         elif dim == 2:
104             return nn.BatchNorm2d(num_channels)
105     elif normalization == 'instance':
106         if dim == 3:
107             return nn.InstanceNorm3d(num_channels)
108         elif dim == 2:
109             return nn.InstanceNorm2d(num_channels)
110     elif 'group' in normalization:
111         num_groups = int(normalization.partition('group')[-1]) # get the group size for
112         return nn.GroupNorm(num_groups=num_groups, num_channels=num_channels)
113
114
115 class Concatenate(nn.Module):
116     def __init__(self):
117         super(Concatenate, self).__init__()
118

```

```

119     def forward(self, layer_1, layer_2):
120         x = torch.cat((layer_1, layer_2), 1)
121
122         return x
123
124
125 class DownBlock(nn.Module):
126     """
127     A helper Module that performs 2 Convolutions and 1 MaxPool.
128     An activation follows each convolution.
129     A normalization layer follows each convolution.
130     """
131
132     def __init__(self,
133                 in_channels: int,
134                 out_channels: int,
135                 pooling: bool = True,
136                 activation: str = 'relu',
137                 normalization: str = None,
138                 dim: str = 2,
139                 conv_mode: str = 'same'):
140         super().__init__()
141
142         self.in_channels = in_channels
143         self.out_channels = out_channels
144         self.pooling = pooling
145         self.normalization = normalization
146         if conv_mode == 'same':
147             self.padding = 1
148         elif conv_mode == 'valid':
149             self.padding = 0
150         self.dim = dim
151         self.activation = activation
152
153         # conv layers
154         self.conv1 = get_conv_layer(self.in_channels, self.out_channels, kernel_size=3,
155                                     bias=True, dim=self.dim)
156         self.conv2 = get_conv_layer(self.out_channels, self.out_channels, kernel_size=3
157                                     bias=True, dim=self.dim)
158
159         # pooling layer
160         if self.pooling:
161             self.pool = get_maxpool_layer(kernel_size=2, stride=2, padding=0, dim=self.
162

```

```

163         # activation layers
164         self.act1 = get_activation(self.activation)
165         self.act2 = get_activation(self.activation)
166
167         # normalization layers
168         if self.normalization:
169             self.norm1 = get_normalization(normalization=self.normalization, num_channels=self.num_channels,
170                                           dim=self.dim)
171             self.norm2 = get_normalization(normalization=self.normalization, num_channels=self.num_channels,
172                                           dim=self.dim)
173
174     def forward(self, x):
175         y = self.conv1(x) # convolution 1
176         y = self.act1(y) # activation 1
177         if self.normalization:
178             y = self.norm1(y) # normalization 1
179         y = self.conv2(y) # convolution 2
180         y = self.act2(y) # activation 2
181         if self.normalization:
182             y = self.norm2(y) # normalization 2
183
184         before_pooling = y # save the outputs before the pooling operation
185         if self.pooling:
186             y = self.pool(y) # pooling
187         return y, before_pooling
188
189
190     class UpBlock(nn.Module):
191         """
192         A helper Module that performs 2 Convolutions and 1 UpConvolution/Upsample.
193         An activation follows each convolution.
194         A normalization layer follows each convolution.
195         """
196
197         def __init__(self,
198                     in_channels: int,
199                     out_channels: int,
200                     activation: str = 'relu',
201                     normalization: str = None,
202                     dim: int = 3,
203                     conv_mode: str = 'same',
204                     up_mode: str = 'transposed'
205                     ):
206             super().__init__()
207

```

```

208         self.in_channels = in_channels
209         self.out_channels = out_channels
210         self.normalization = normalization
211         if conv_mode == 'same':
212             self.padding = 1
213         elif conv_mode == 'valid':
214             self.padding = 0
215         self.dim = dim
216         self.activation = activation
217         self.up_mode = up_mode
218
219         # upconvolution/upsample layer
220         self.up = get_up_layer(self.in_channels, self.out_channels, kernel_size=2, stri
221                                up_mode=self.up_mode)
222
223         # conv layers
224         self.conv0 = get_conv_layer(self.in_channels, self.out_channels, kernel_size=1,
225                                     bias=True, dim=self.dim)
226         self.conv1 = get_conv_layer(2 * self.out_channels, self.out_channels, kernel_si
227                                     padding=self.padding,
228                                     bias=True, dim=self.dim)
229         self.conv2 = get_conv_layer(self.out_channels, self.out_channels, kernel_size=3
230                                     bias=True, dim=self.dim)
231
232         # activation layers
233         self.act0 = get_activation(self.activation)
234         self.act1 = get_activation(self.activation)
235         self.act2 = get_activation(self.activation)
236
237         # normalization layers
238         if self.normalization:
239             self.norm0 = get_normalization(normalization=self.normalization, num_channe
240                                             dim=self.dim)
241             self.norm1 = get_normalization(normalization=self.normalization, num_channe
242                                             dim=self.dim)
243             self.norm2 = get_normalization(normalization=self.normalization, num_channe
244                                             dim=self.dim)
245
246         # concatenate layer
247         self.concat = Concatenate()
248
249     def forward(self, encoder_layer, decoder_layer):
250         """ Forward pass
251         Arguments:
252             encoder_layer: Tensor from the encoder pathway

```



```

252         encoder_layer: Tensor from the encoder pathway
253         decoder_layer: Tensor from the decoder pathway (to be up'd)
254     """
255     up_layer = self.up(decoder_layer) # up-convolution/up-sampling
256     cropped_encoder_layer, dec_layer = autocrop(encoder_layer, up_layer) # cropping
257
258     if self.up_mode != 'transposed':
259         # We need to reduce the channel dimension with a conv layer
260         up_layer = self.conv0(up_layer) # convolution 0
261     up_layer = self.act0(up_layer) # activation 0
262     if self.normalization:
263         up_layer = self.norm0(up_layer) # normalization 0
264
265     merged_layer = self.concat(up_layer, cropped_encoder_layer) # concatenation
266     y = self.conv1(merged_layer) # convolution 1
267     y = self.act1(y) # activation 1
268     if self.normalization:
269         y = self.norm1(y) # normalization 1
270     y = self.conv2(y) # convolution 2
271     y = self.act2(y) # activation 2
272     if self.normalization:
273         y = self.norm2(y) # normalization 2
274     return y
275
276
277 class UNet(nn.Module):
278     def __init__(self,
279                 in_channels: int = 1,
280                 out_channels: int = 2,
281                 n_blocks: int = 4,
282                 start_filters: int = 32,
283                 activation: str = 'relu',
284                 normalization: str = 'batch',
285                 conv_mode: str = 'same',
286                 dim: int = 2,
287                 up_mode: str = 'transposed'
288             ):
289         super().__init__()
290
291         self.in_channels = in_channels
292         self.out_channels = out_channels
293         self.n_blocks = n_blocks
294         self.start_filters = start_filters
295         self.activation = activation
296         self.normalization = normalization

```

```

297         self.conv_mode = conv_mode
298         self.dim = dim
299         self.up_mode = up_mode
300
301         self.down_blocks = []
302         self.up_blocks = []
303
304         # create encoder path
305         for i in range(self.n_blocks):
306             num_filters_in = self.in_channels if i == 0 else num_filters_out
307             num_filters_out = self.start_filters * (2 ** i)
308             pooling = True if i < self.n_blocks - 1 else False
309
310             down_block = DownBlock(in_channels=num_filters_in,
311                                   out_channels=num_filters_out,
312                                   pooling=pooling,
313                                   activation=self.activation,
314                                   normalization=self.normalization,
315                                   conv_mode=self.conv_mode,
316                                   dim=self.dim)
317
318             self.down_blocks.append(down_block)
319
320         # create decoder path (requires only n_blocks-1 blocks)
321         for i in range(n_blocks - 1):
322             num_filters_in = num_filters_out
323             num_filters_out = num_filters_in // 2
324
325             up_block = UpBlock(in_channels=num_filters_in,
326                               out_channels=num_filters_out,
327                               activation=self.activation,
328                               normalization=self.normalization,
329                               conv_mode=self.conv_mode,
330                               dim=self.dim,
331                               up_mode=self.up_mode)
332
333             self.up_blocks.append(up_block)
334
335         # final convolution
336         self.conv_final = get_conv_layer(num_filters_out, self.out_channels, kernel_size,
337                                          bias=True, dim=self.dim)
338
339         # add the list of modules to current module
340         self.down_blocks = nn.ModuleList(self.down_blocks)
341         self.up_blocks = nn.ModuleList(self.up_blocks)

```

```

342
343     # initialize the weights
344     self.initialize_parameters()
345
346     @staticmethod
347     def weight_init(module, method, **kwargs):
348         if isinstance(module, (nn.Conv3d, nn.Conv2d, nn.ConvTranspose3d, nn.ConvTranspose2d)):
349             method(module.weight, **kwargs) # weights
350
351     @staticmethod
352     def bias_init(module, method, **kwargs):
353         if isinstance(module, (nn.Conv3d, nn.Conv2d, nn.ConvTranspose3d, nn.ConvTranspose2d)):
354             method(module.bias, **kwargs) # bias
355
356     def initialize_parameters(self,
357                             method_weights=nn.init.xavier_uniform_,
358                             method_bias=nn.init.zeros_,
359                             kwargs_weights={},
360                             kwargs_bias={}
361                             ):
362         for module in self.modules():
363             self.weight_init(module, method_weights, **kwargs_weights) # initialize weights
364             self.bias_init(module, method_bias, **kwargs_bias) # initialize bias
365
366     def forward(self, x: torch.tensor):
367         encoder_output = []
368
369         # Encoder pathway
370         for module in self.down_blocks:
371             x, before_pooling = module(x)
372             encoder_output.append(before_pooling)
373
374         # Decoder pathway
375         for i, module in enumerate(self.up_blocks):
376             before_pool = encoder_output[-(i + 2)]
377             x = module(before_pool, x)
378
379         x = self.conv_final(x)
380
381         return x

```

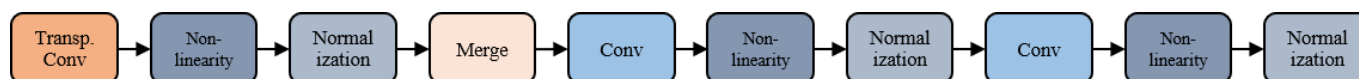
path. As you can see in `unet.py` the `DownBlock` and the `UpBlock` help to build the architecture. Both use smaller helper functions that return the correct layer, depending on what arguments are passed, e.g. if a 2D (`dim=2`) or 3D (`dim=3`) network is wanted. The number of blocks is defined by the `depth` of the network.

A `DownBlock` generally has the following scheme:



DownBlock — Image by Johannes Schmidt

A `UpBlock` has the following layers:



UpBlock — Image by Johannes Schmidt

For our `Unet` class we just need to combine these blocks and make sure that the correct layers from the encoder are concatenated to the decoder (skip pathways). These layers have to be cropped if their sizes do not match with the corresponding layers from the decoder. In such cases, the `autocrop` function is used. For merging, I concatenate along the channel dimension (see `Concatenate`). Instead of transposed convolutions we could also use upsampling layers (interpolation methods) that are followed by a `1x1` or `3x3` convolution block to reduce the channel dimension. Using interpolation generally gets rid of the checkerboard artifact. For 3D input consider using trilinear interpolation.

At the end we just need to think about the parameter initialization. By default, the weights are initialized with `torch.nn.init.xavier_uniform_` and the biases are initialized with zeros using `torch.nn.init.zeros_`.

For details and the available parameter options, I encourage you to take a look at the code. Feel free to change the code to your needs or expand e.g. the number of activation functions.

Creating a U-Net model

Let's create such a model and use it to make a prediction on some random input:

```

from unet import UNet

model = UNet(in_channels=1,
             out_channels=2,
             n_blocks=4,
             start_filters=32,
             activation='relu',
             normalization='batch',
             conv_mode='same',
             dim=2)

x = torch.randn(size=(1, 1, 512, 512), dtype=torch.float32)
with torch.no_grad():
    out = model(x)

print(f'Out: {out.shape}')
```

This will give us:

```
Out: torch.Size([1, 2, 512, 512])
```

To check whether our model is correct, we can get the model's summary with this package [pytorch-summary](#):

```

from torchsummary import summary
summary = summary(model, (1, 512, 512))
```

which prints out a summary like this:

1	=====		
2	Layer (type:depth-idx)	Output Shape	Param #
3	=====		
4	└─ModuleList: 1	[]	--
5	└─DownBlock: 2-1	[-1, 32, 256, 256]	--
6	└─Conv2d: 3-1	[-1, 32, 512, 512]	320
7	└─ReLU: 3-2	[-1, 32, 512, 512]	--
8	└─BatchNorm2d: 3-3	[-1, 32, 512, 512]	64
9	└─Conv2d: 3-4	[-1, 32, 512, 512]	9,248
10	└─ReLU: 3-5	[-1, 32, 512, 512]	--
11	└─BatchNorm2d: 3-6	[-1, 32, 512, 512]	64

11			└─BatchNorm2d: 3-6	[-1, 32, 32, 32]	0
12			└─MaxPool2d: 3-7	[-1, 32, 256, 256]	--
13			└─DownBlock: 2-2	[-1, 64, 128, 128]	--
14			└─Conv2d: 3-8	[-1, 64, 256, 256]	18,496
15			└─ReLU: 3-9	[-1, 64, 256, 256]	--
16			└─BatchNorm2d: 3-10	[-1, 64, 256, 256]	128
17			└─Conv2d: 3-11	[-1, 64, 256, 256]	36,928
18			└─ReLU: 3-12	[-1, 64, 256, 256]	--
19			└─BatchNorm2d: 3-13	[-1, 64, 256, 256]	128
20			└─MaxPool2d: 3-14	[-1, 64, 128, 128]	--
21			└─DownBlock: 2-3	[-1, 128, 64, 64]	--
22			└─Conv2d: 3-15	[-1, 128, 128, 128]	73,856
23			└─ReLU: 3-16	[-1, 128, 128, 128]	--
24			└─BatchNorm2d: 3-17	[-1, 128, 128, 128]	256
25			└─Conv2d: 3-18	[-1, 128, 128, 128]	147,584
26			└─ReLU: 3-19	[-1, 128, 128, 128]	--
27			└─BatchNorm2d: 3-20	[-1, 128, 128, 128]	256
28			└─MaxPool2d: 3-21	[-1, 128, 64, 64]	--
29			└─DownBlock: 2-4	[-1, 256, 64, 64]	--
30			└─Conv2d: 3-22	[-1, 256, 64, 64]	295,168
31			└─ReLU: 3-23	[-1, 256, 64, 64]	--
32			└─BatchNorm2d: 3-24	[-1, 256, 64, 64]	512
33			└─Conv2d: 3-25	[-1, 256, 64, 64]	590,080
34			└─ReLU: 3-26	[-1, 256, 64, 64]	--
35			└─BatchNorm2d: 3-27	[-1, 256, 64, 64]	512
36			└─ModuleList: 1	[]	--
37			└─UpBlock: 2-5	[-1, 128, 128, 128]	--
38			└─ConvTranspose2d: 3-28	[-1, 128, 128, 128]	131,200
39			└─ReLU: 3-29	[-1, 128, 128, 128]	--
40			└─BatchNorm2d: 3-30	[-1, 128, 128, 128]	256
41			└─Concatenate: 3-31	[-1, 256, 128, 128]	--
42			└─Conv2d: 3-32	[-1, 128, 128, 128]	295,040
43			└─ReLU: 3-33	[-1, 128, 128, 128]	--
44			└─BatchNorm2d: 3-34	[-1, 128, 128, 128]	256
45			└─Conv2d: 3-35	[-1, 128, 128, 128]	147,584
46			└─ReLU: 3-36	[-1, 128, 128, 128]	--
47			└─BatchNorm2d: 3-37	[-1, 128, 128, 128]	256
48			└─UpBlock: 2-6	[-1, 64, 256, 256]	--
49			└─ConvTranspose2d: 3-38	[-1, 64, 256, 256]	32,832
50			└─ReLU: 3-39	[-1, 64, 256, 256]	--
51			└─BatchNorm2d: 3-40	[-1, 64, 256, 256]	128
52			└─Concatenate: 3-41	[-1, 128, 256, 256]	--
53			└─Conv2d: 3-42	[-1, 64, 256, 256]	73,792
54			└─ReLU: 3-43	[-1, 64, 256, 256]	--
55			└─BatchNorm2d: 3-44	[-1, 64, 256, 256]	128

56			└─Conv2d: 3-45	[-1, 64, 256, 256]	36,928
57			└─ReLU: 3-46	[-1, 64, 256, 256]	--
58			└─BatchNorm2d: 3-47	[-1, 64, 256, 256]	128
59			└─UpBlock: 2-7	[-1, 32, 512, 512]	--
60			└─ConvTranspose2d: 3-48	[-1, 32, 512, 512]	8,224
61			└─ReLU: 3-49	[-1, 32, 512, 512]	--
62			└─BatchNorm2d: 3-50	[-1, 32, 512, 512]	64
63			└─Concatenate: 3-51	[-1, 64, 512, 512]	--
64			└─Conv2d: 3-52	[-1, 32, 512, 512]	18,464
65			└─ReLU: 3-53	[-1, 32, 512, 512]	--
66			└─BatchNorm2d: 3-54	[-1, 32, 512, 512]	64
67			└─Conv2d: 3-55	[-1, 32, 512, 512]	9,248
68			└─ReLU: 3-56	[-1, 32, 512, 512]	--
69			└─BatchNorm2d: 3-57	[-1, 32, 512, 512]	64
70			└─Conv2d: 1-1	[-1, 2, 512, 512]	66
71	=====				
72	Total params: 1,928,322				
73	Trainable params: 1,928,322				
74	Non-trainable params: 0				
75	Total mult-adds (M): -14267.05				
76	=====				
77	Input size (MB): 1.00				
78	Forward/backward pass size (MB): 1156.00				
79	Params size (MB): 7.36				
80	Estimated Total Size (MB): 1164.36				
81	=====				

About input sizes

To ensure correct semantic concatenations, it is advised to use input sizes that return even spatial dimensions in every block but the last in the encoder. For example: An input size of 120^2 gives intermediate output shapes of $[60^2, 30^2, 15^2]$ in the encoder path for a U-Net with `depth=4`. A U-Net with `depth=5` with the same input size is not recommended, as a maxpooling operation on odd spatial dimensions (e.g. on a 15^2 input) should be avoided.

To make our lives easier, we can numerically compute the maximum network depth for a given input dimension with a simple function:

```
shape = 1920
```

```
def compute_max_depth(shape, max_depth=10, print_out=True):
    shapes = []
    shapes.append(shape)
    for level in range(1, max_depth):
        if shape % 2 ** level == 0 and shape / 2 ** level > 1:
            shapes.append(shape / 2 ** level)
            if print_out:
                print(f'Level {level}: {shape / 2 ** level}')
        else:
            if print_out:
                print(f'Max-level: {level - 1}')
            break

    return shapes

out = compute_max_depth(shape, print_out=True, max_depth=10)
```

This will output

```
Level 1: 960.0
Level 2: 480.0
Level 3: 240.0
Level 4: 120.0
Level 5: 60.0
Level 6: 30.0
Level 7: 15.0
Max-level: 7
```

which tells us that that we can design a U-Net as deep as this without having to worry about semantic mismatches. Conversely, we can also numerically determine the possible input shapes dimensions for a given depth:

```
low = 10
high = 512
depth = 8

def compute_possible_shapes(low, high, depth):
    possible_shapes = {}
    for shape in range(low, high + 1):
        shapes = compute_max_depth(shape,
                                    max_depth=depth,
                                    print_out=False)
```



```
        if len(shapes) == depth:
            possible_shapes[shape] = shapes

    return possible_shapes

possible_shapes = compute_possible_shapes(low, high, depth)
```

This will output

```
{256: [256, 128.0, 64.0, 32.0, 16.0, 8.0, 4.0, 2.0],
384: [384, 192.0, 96.0, 48.0, 24.0, 12.0, 6.0, 3.0],
512: [512, 256.0, 128.0, 64.0, 32.0, 16.0, 8.0, 4.0]}
```

which tells us that we can have 3 different input shapes with such a level 8 U-Net architecture. But I dare to say that such a network with this input size is probably not useful in practice.

Summary

In this part we created a configurable UNet model for the purpose of semantic segmentation. Now that we have built our model, it is time to create a training loop in the [next chapter](#).

Sign up for The Variable

By Towards Data Science

Every Thursday, the Variable delivers the very best of Towards Data Science: from hands-on tutorials and cutting-edge research to original features you don't want to miss. [Take a look.](#)

Get this newsletter

Emails will be sent to mr.spons@gmail.com.
[Not you?](#)

[About](#) [Help](#) [Legal](#)

Get the Medium app

