



# Introduction to **Machine Learning and Data Mining** (Học máy và Khai phá dữ liệu)

**Khoat Than**

School of Information and Communication Technology  
Hanoi University of Science and Technology

2022


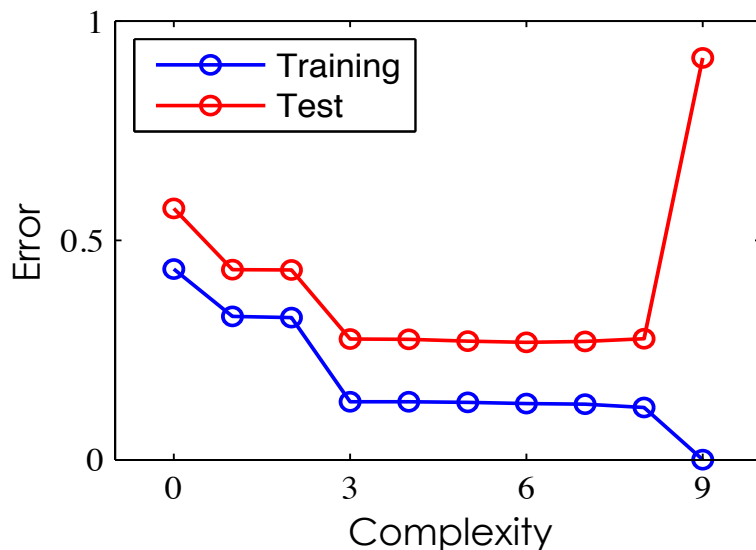
# Contents

---

- Introduction to Machine Learning & Data Mining
- Supervised learning
- Unsupervised learning
- Probabilistic modeling
- **Regularization**
- Practical advice

# Revisiting overfitting

- The complexity of the learned function:  $y = \hat{f}(x, \mathbf{D})$ 
  - For a given training data  $\mathbf{D}$ : the more complicated  $\hat{f}$ , the more possibility that  $\hat{f}$  fits  $\mathbf{D}$  better.
  - For a given  $\mathbf{D}$ : there exist many functions that fit  $\mathbf{D}$  perfectly (i.e., no error on  $\mathbf{D}$ ).
  - However, those functions might generalize badly.



# The Bias-Variance Decomposition

- Consider  $y = f(x) + \epsilon$  as the regression function
  - ❖ where  $\epsilon \sim \mathcal{N}(0, \sigma^2)$  is a Gaussian noise with mean 0 and variance  $\sigma^2$ .
  - ❖  $\epsilon$  may represent the *noise* due to measurement or data collection.
- Let  $\hat{f}(x; \mathbf{D})$  be the regressor learned from a training data  $\mathbf{D}$
- Note:
  - ❖ We want that  $\hat{f}(x; \mathbf{D})$  approximates the truth  $f(x)$  well.
  - ❖  $\hat{f}(x; \mathbf{D})$  is random, according to the randomness when collecting  $\mathbf{D}$ .
- For any  $x$ , the error made by  $\hat{f}(x; \mathbf{D})$  is

$$\mathbb{E}_{D, \epsilon} \left( y(x) - \hat{f}(x; \mathbf{D}) \right)^2 = \sigma^2 + \text{Bias}^2 \left( \hat{f}(x; \mathbf{D}) \right) + \text{Var} \left( \hat{f}(x; \mathbf{D}) \right)$$

$$\text{Bias} \left( \hat{f}(x; \mathbf{D}) \right) = \mathbb{E}_D [f(x) - \hat{f}(x; \mathbf{D})]$$

$$\text{Var} \left( \hat{f}(x; \mathbf{D}) \right) = \mathbb{E}_D \left( \hat{f}(x; \mathbf{D}) - \mathbb{E}_D \hat{f}(x; \mathbf{D}) \right)^2$$


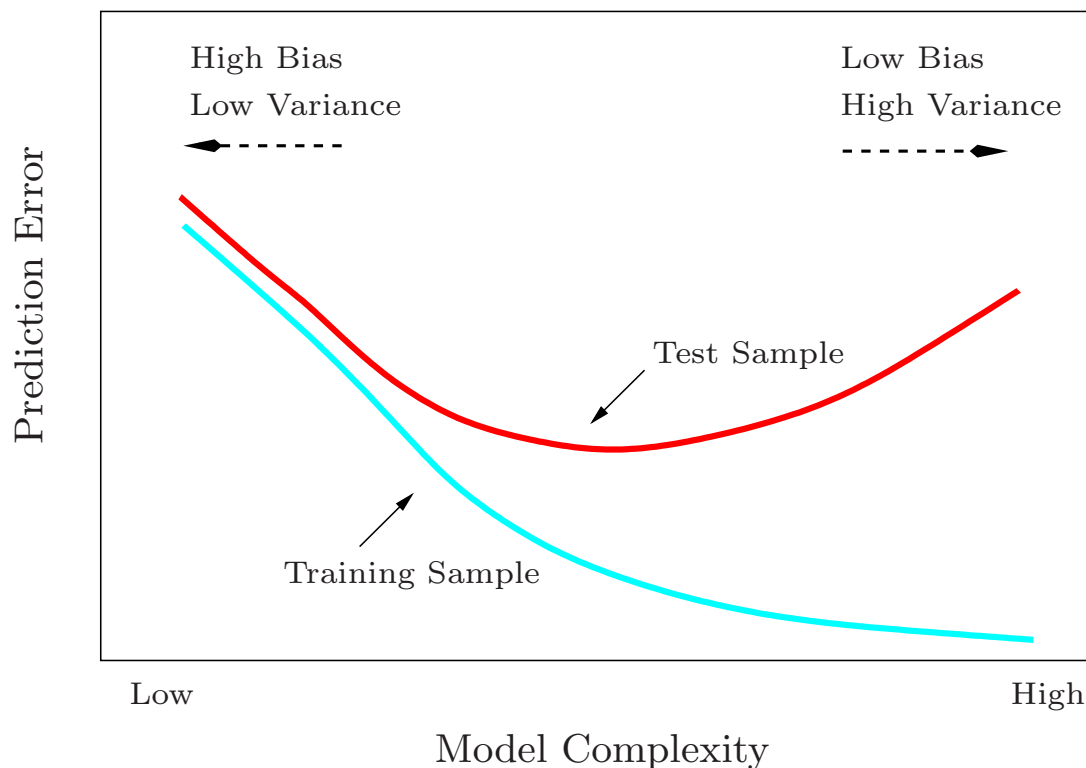
# The Bias-Variance Decomposition (2)

$$\begin{aligned} \text{Error}(x) &= \sigma^2 + \text{Bias}^2 \left( \hat{f}(x; \mathbf{D}) \right) + \text{Var} \left( \hat{f}(x; \mathbf{D}) \right) \\ &= \text{Irreducible Error} + \text{Bias}^2 + \text{Variance} \end{aligned}$$

- This is known as the **Bias-Variance Decomposition**
  - ❖ *Irreducible Error*: cannot be avoided due to noises or uncontrolled factors
  - ❖ *Bias*: the average of our estimate differs from the true mean
  - ❖ *Variance*: the expected squared deviation of  $\hat{f}(x; \mathbf{D})$  around its mean


# Bias-Variance tradeoff: classical view

- The more complex the model  $\hat{f}(x; \mathbf{D})$  is, the more data points it can capture, and the lower the bias can be.
- ❖ However, higher complexity will make the model "move" more to capture the data points, and hence its variance will be larger.




# Regularization: introduction

- *Regularization* is now a popular and useful technique in ML.
- It is a technique to exploit further information to
  - Reduce overfitting in ML.
  - Solve ill-posed problems in Maths.
- The further information is often enclosed in a *penalty on the complexity* of  $\hat{f}(x, \mathbf{D})$ .
  - More penalty will be imposed on complex functions.
  - We prefer simpler functions among all that fit well the training data.



Tikhonov,  
smoothing an ill-  
posed problem



Zarembka, model  
complexity  
minimization



Bayes: priors  
over parameters



Andrew Ng: need no  
maths, but it prevents  
overfitting!

# Regularization in Ridge regression

- Learning a linear regressor by ordinary least squares (OLS) from a training data  $\mathbf{D} = \{(x_1, y_1), \dots, (x_M, y_M)\}$  is reduced to the following problem:

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2 = \arg \min_w \sum_{(x_i, y_i) \in \mathbf{D}} (y_i - w^T x_i)^2$$

- For Ridge regression, learning is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2$$

- Where  $\lambda$  is a positive constant.
- The term  $\lambda \|w\|_2^2$  plays the role as *limiting the size/complexity of  $w$* .
- $\lambda$  allows us to trade off between fitness on  $\mathbf{D}$  and generalization on future observations.
- Ridge regression is a regularized version of OLS.



# Regularization: the principle

- We need to learn a function  $f(x, w)$  from the training set  $\mathbf{D}$ 
  - $x$  is a data example and belongs to **input space**.
  - $w$  is the parameter and often belongs to a **parameter space**  $\mathbf{W}$ .
  - $\mathbf{F} = \{f(x, w) : w \in \mathbf{W}\}$  is the **function space**, parameterized by  $w$ .
- For many ML models, the training problem is often reduced to the following optimization:

$$w^* = \arg \min_{w \in \mathbf{W}} L(f(x, w), \mathbf{D}) \quad (1)$$

- $w$  sometimes tells the size/complexity of that function.
  - $L(f(x, w), \mathbf{D})$  is an **empirical loss/risk** which depends on  $\mathbf{D}$ . This loss shows how well function  $f$  fits  $\mathbf{D}$ .
- Another view:  $f^* = \arg \min_{f \in \mathbf{F}} L(f(x, w), \mathbf{D})$

# Regularization: the principle

- Adding a penalty to (1), we consider

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w) \quad (2)$$

- Where  $\lambda > 0$  is called *the regularization/penalty constant*.
- $g(w)$  measures the complexity of  $w$ . ( $g(w) \geq 0$ )
- $L(f(x, w), \mathbf{D})$  measures the goodness of function  $f$  on  $\mathbf{D}$ .
- The penalty (regularization) term:  $\lambda g(w)$ 
  - Allows to trade off the fitness on  $\mathbf{D}$  and the generalization.
  - The greater  $\lambda$ , the heavier penalty, implying that  $g(w)$  should be smaller.
  - In practice,  $\lambda$  should be neither too small nor too large.  
( $\lambda$  không nên quá lớn hoặc quá bé trong thực tế)

# Regularization: popular types



- $g(w)$  often relates to some norms when  $w$  is an  $n$ -dimensional vector.

□  $L_0$ -norm:  $\|w\|_0$  counts the number of non-zeros in  $w$ .

□  $L_1$ -norm: 
$$\|w\|_1 = \sum_{i=1}^n |w_i|$$

□  $L_2$ -norm: 
$$\|w\|_2^2 = \sum_{i=1}^n w_i^2$$

□  $L_p$ -norm: 
$$\|w\|_p = \sqrt[p]{|w_1|^p + \dots + |w_n|^p}$$



# Regularization in Ridge regression

- Ridge regression can be derived from OLS by adding a penalty term into the objective function when learning.
- Learning a regressor in Ridge is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_2^2$$

- Where  $\lambda$  is a positive constant.
- The term  $\lambda \|w\|_2^2$  plays the role as regularization.
- Large  $\lambda$  reduces the size of  $w$ .

# Regularization in Lasso

- Lasso [Tibshirani, 1996] is a variant of OLS for linear regression by using  $L_1$  to do regularization.
- Learning a linear regressor is reduced to

$$w^* = \arg \min_w RSS(w, \mathbf{D}) + \lambda \|w\|_1$$

- Where  $\lambda$  is a positive constant.
- $\lambda \|w\|_1$  is the regularization term. Large  $\lambda$  reduces the size of  $w$ .
- Regularization here amounts to imposing a Laplace distribution (as prior) over each  $w_i$ , with density function:

$$p(w_i | \lambda) = \frac{\lambda}{2} e^{-\lambda |w_i|}$$

- The larger  $\lambda$ , the more possibility that  $w_i = 0$ .

# Regularization in SVM

- Learning a classifier in SVM is reduced to the following problem:

- Minimize 
$$\frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2}$$

- Conditioned on  $y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1, \quad \forall i = 1..r$

- In the cases of noises/errors, learning is reduced to

- Minimize 
$$\frac{\langle \mathbf{w} \cdot \mathbf{w} \rangle}{2} + C \sum_{i=1}^r \xi_i$$

- Conditioned on 
$$\begin{cases} y_i(\langle \mathbf{w} \cdot \mathbf{x}_i \rangle + b) \geq 1 - \xi_i, & \forall i = 1..r \\ \xi_i \geq 0, & \forall i = 1..r \end{cases}$$

- $C(\xi_1 + \dots + \xi_r)$  is *the regularization term*.

# Some other regularizations

---

- **Dropout:** (by Hilton and his colleagues, 2012)

- At each iteration of the training process, randomly drop out some parts and just update the other parts of our model.

- **Batch normalization** [Ioffe & Szegedy, 2015]

- Normalize the inputs at each neuron of a neural network
- Reduce input variance, easier training, faster convergence

- **Data augmentation**

- Produce different versions of an example in the training set, by adding simple noises, translation, rotation, cropping, ...
- Those versions are added to the training data set

- **Early stopping**

- Stop training early to avoid overtraining & reduce overfitting

# Regularization: MAP role

- Under some conditions, we can view regularization as

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w)$$

Likelihood

Prior

- Where  $\mathbf{D}$  is a sample from a probability distribution whose log likelihood is  $-L(f(x, w), \mathbf{D})$ .

- $w$  is a random variable and follows the prior with density  
 $p(w) \propto \exp(-\lambda g(w))$

- Then  $w^* = \arg \max_{w \in \mathcal{W}} \{-L(f(x, w), \mathbf{D}) - \lambda g(w)\}$

$$w^* = \arg \max_{w \in \mathcal{W}} \log \Pr(\mathbf{D}|w) + \log \Pr(w) = \arg \max_{w \in \mathcal{W}} \log \Pr(w|\mathbf{D})$$

- As a result, regularization in fact helps us to learn an MAP solution  $w^*$ .



# Regularization: MAP in Ridge

- Consider the Gaussian regression model:

- $w$  follows a Gaussian prior:  $N(w \mid 0, \sigma^2 \rho^2)$ .
- Variable  $f = y - w^T x$  follows the Gaussian distribution  $N(f \mid 0, \rho^2, w)$  with mean 0 and variance  $\rho^2$ , and conditioned on  $w$ .

- Then the MAP estimation of  $f$  from the training data  $\mathbf{D}$  is

$$w^* = \operatorname{argmax}_w \log \Pr(w \mid D) = \operatorname{argmax}_w \log [\Pr(D \mid w) * \Pr(w)]$$

$$= \operatorname{argmin}_w \sum_{(x_i, y_i)} \frac{1}{2\rho^2} (y_i - w^T x_i)^2 + \frac{1}{2\sigma^2 \rho^2} w^T w - \text{constant}$$

$$= \operatorname{argmin}_w \sum_{(x_i, y_i)} (y_i - w^T x_i)^2 + \frac{1}{\sigma^2} w^T w$$

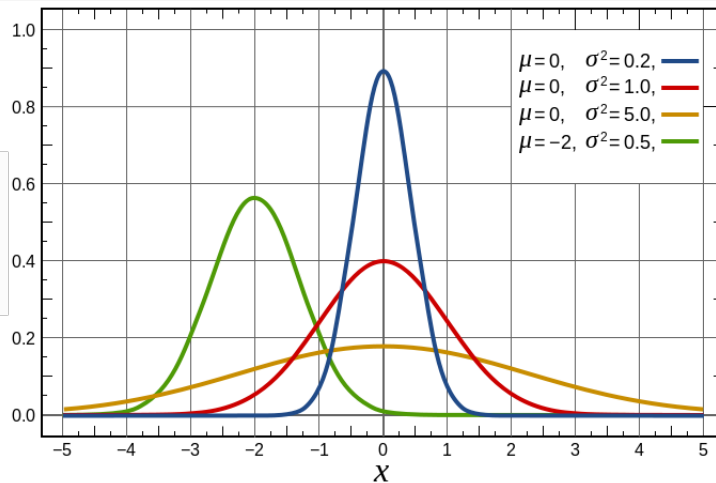
**Ridge regression**  
@@

- Regularization using  $L_2$  with penalty constant  $\lambda = \sigma^2$ .

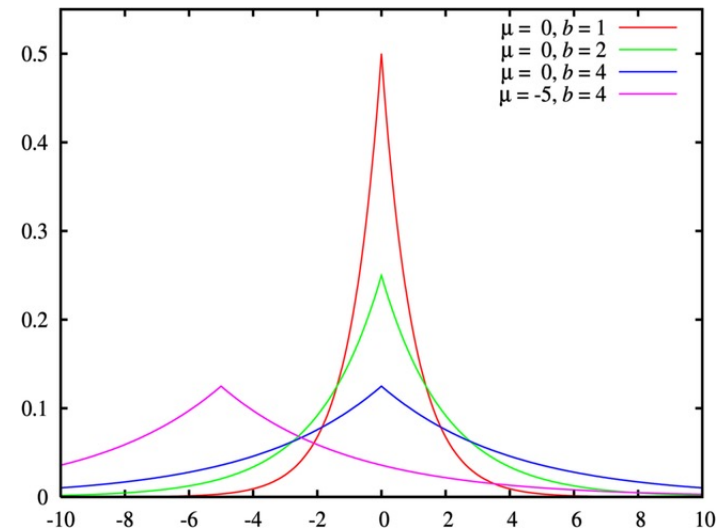
# Regularization: MAP in Ridge & Lasso

- The regularization constant in Ridge:  $\lambda = \sigma^{-2}$
- The regularization constant in Lasso:  $\lambda = b^{-1}$
- Gaussian (left) and Laplace distribution (right)

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$




$$f(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right)$$




# Regularization: limiting the search space

- The regularization constant in Ridge:  $\lambda = \sigma^{-2}$
- The regularization constant in Lasso:  $\lambda = b^{-1}$
- *The larger  $\lambda$ , the higher probability that  $x$  occurs around 0.*

$$f(x, \mu, \sigma) = \frac{1}{\sigma\sqrt{2\pi}} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$



$$f(x|\mu, b) = \frac{1}{2b} \exp\left(-\frac{|x - \mu|}{b}\right)$$



# Regularization: limiting the search space

- The regularized problem:

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) + \lambda g(w) \quad (2)$$


- A result from the optimization literature shows that (2) is equivalent to the following:

$$w^* = \arg \min_{w \in \mathcal{W}} L(f(x, w), \mathbf{D}) \quad \text{such that} \quad g(w) \leq s \quad (3)$$

- For some constant  $s$ .
- *Note that the constraint of  $g(w) \leq s$  plays the role as limiting the search space of  $w$ .*

# Regularization: effects of $\lambda$

- Vector  $\mathbf{w}^* = (w_0, s1, s2, s3, s4, s5, s6, \text{Age}, \text{Sex}, \text{BMI}, \text{BP})$  changes when  $\lambda$  changes in Ridge regression.
  - $\mathbf{w}^*$  goes to 0 as  $\lambda$  increases.



- Ridge regression was under investigation on a prostate dataset with 67 observations.
  - Performance was measured by RMSE (root mean square errors) and Correlation coefficient.

$\lambda$	0.1	1	10	100	1000	10000
RMSE	<b>0.74</b>	<b>0.74</b>	<b>0.74</b>	0.84	1.08	1.16
Correlation coefficient	0.77	0.77	<b>0.78</b>	0.76	0.74	0.73

- Too high or too low values of  $\lambda$  often result in bad predictions.
- Why??

# Bias-Variance tradeoff: revisit

## ■ Classical view:

More complex model  $\hat{f}(x; \mathbf{D})$

- ❖ Lower bias, higher variance

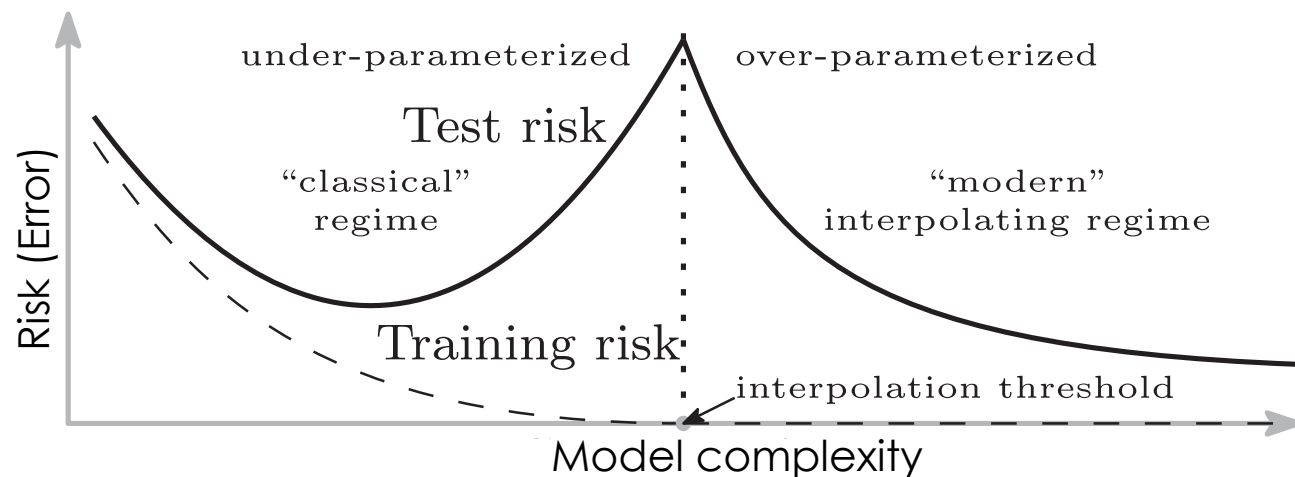

## ■ Modern phenomenon:

- ❖ Very rich models such as neural networks are trained to **exactly fit** the data, but often obtain **high accuracy** on test data [Belkin et al., 2019; Zhang et al., 2021]

- ❖  $\text{Bias} \cong 0$

- ❖ GPT-3, ResNets, VGG, StyleGAN, DALLE-3, ...

## ■ Why???



# Regularization: summary

---

## ■ Advantages:

- Avoid overfitting.
- Limit the search space of the function to be learned.
- Reduce bad effects from noises or errors in observations.
- Might model data better. As an example,  $L_1$  often work well with data/model which are inherently sparse.

## ■ Limitations:

- Consume time to select a good regularization constant.
- Might pose some difficulties to design an efficient algorithm.



# References

---

- Belkin, M., Hsu, D., Ma, S., & Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias–variance trade-off. *Proceedings of the National Academy of Sciences*, 116(32), 15849-15854.
- Ioffe, S., & Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. In *International Conference on Machine Learning* (pp. 448-456).
- Krizhevsky, A., Sutskever, I., & Hinton, G. E. (2012). Imagenet classification with deep convolutional neural networks. *Advances in Neural Information Processing Systems*, 25, 1097-1105.
- Tibshirani, R. (1996). *Regression shrinkage and selection via the Lasso*. *Journal of the Royal Statistical Society*, vol. 58(1), pp. 267-288.
- Trevor Hastie, Robert Tibshirani, Jerome Friedman. *The Elements of Statistical Learning*. Springer, 2009.
- Zhang, C., Bengio, S., Hardt, M., Recht, B., & Vinyals, O. (2021). Understanding deep learning (still) requires rethinking generalization. *Communications of the ACM*, 64(3), 107-115.